

Hands on Kubernetes Container Orchestration In Cluster and Grid Computing



kubernetes

By: Amir Hossein Sorouri
Cluster & Grid Computing
Instructor: Dr. Mohsen Sharifi



IU ST

Date: 2019-May

1. Cluster Review
2. Grid Review
3. Historical Review on Deployment
4. Historical Review on Softwares Architectures

Packed up for the Journey...



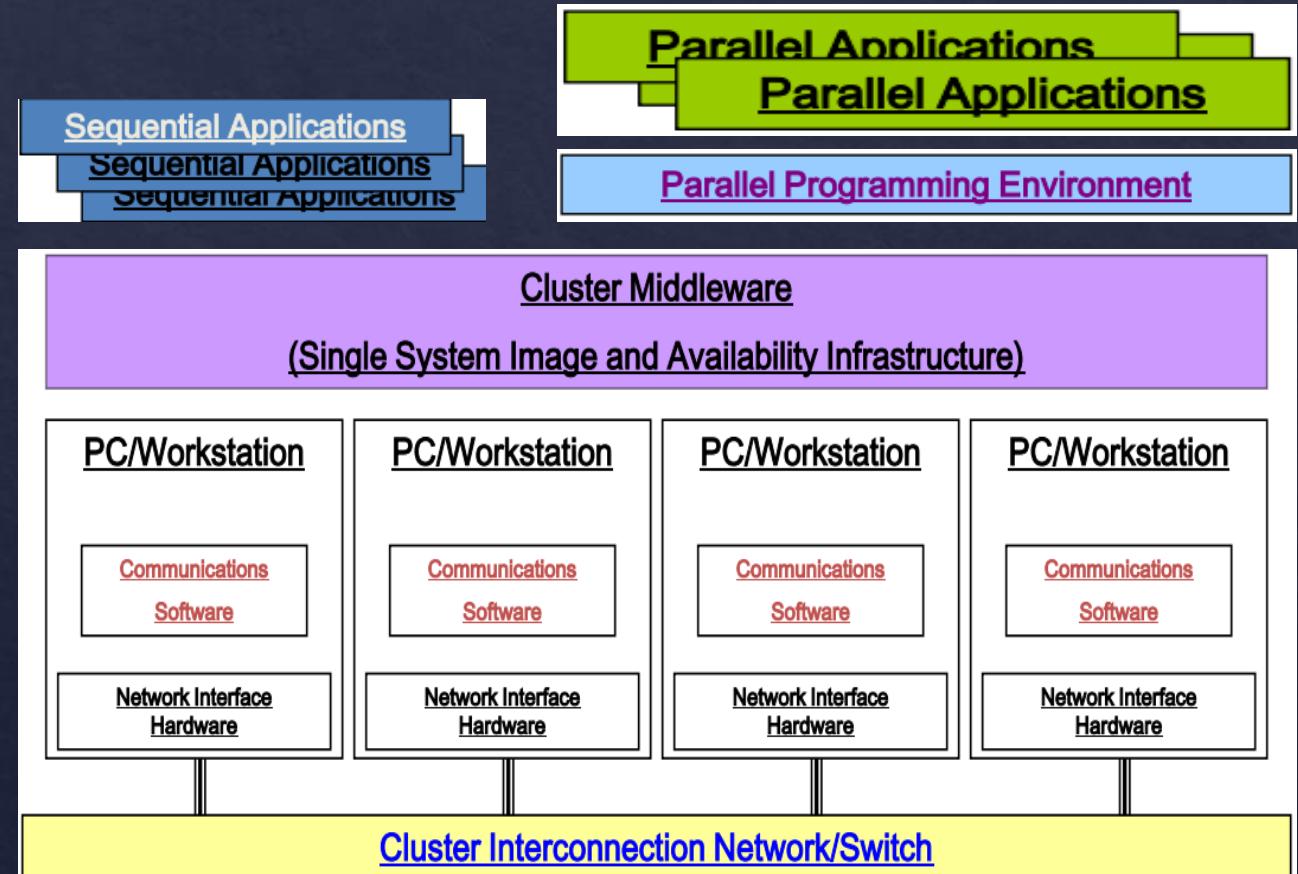
5. Container Concepts and Fundamentals
6. Hands On Docker
7. Hands On Kubernetes



■ Main Issues in Cluster Environment

- Resource Management Scheduling
- Fault Handling
- Migration
- Load Balancing

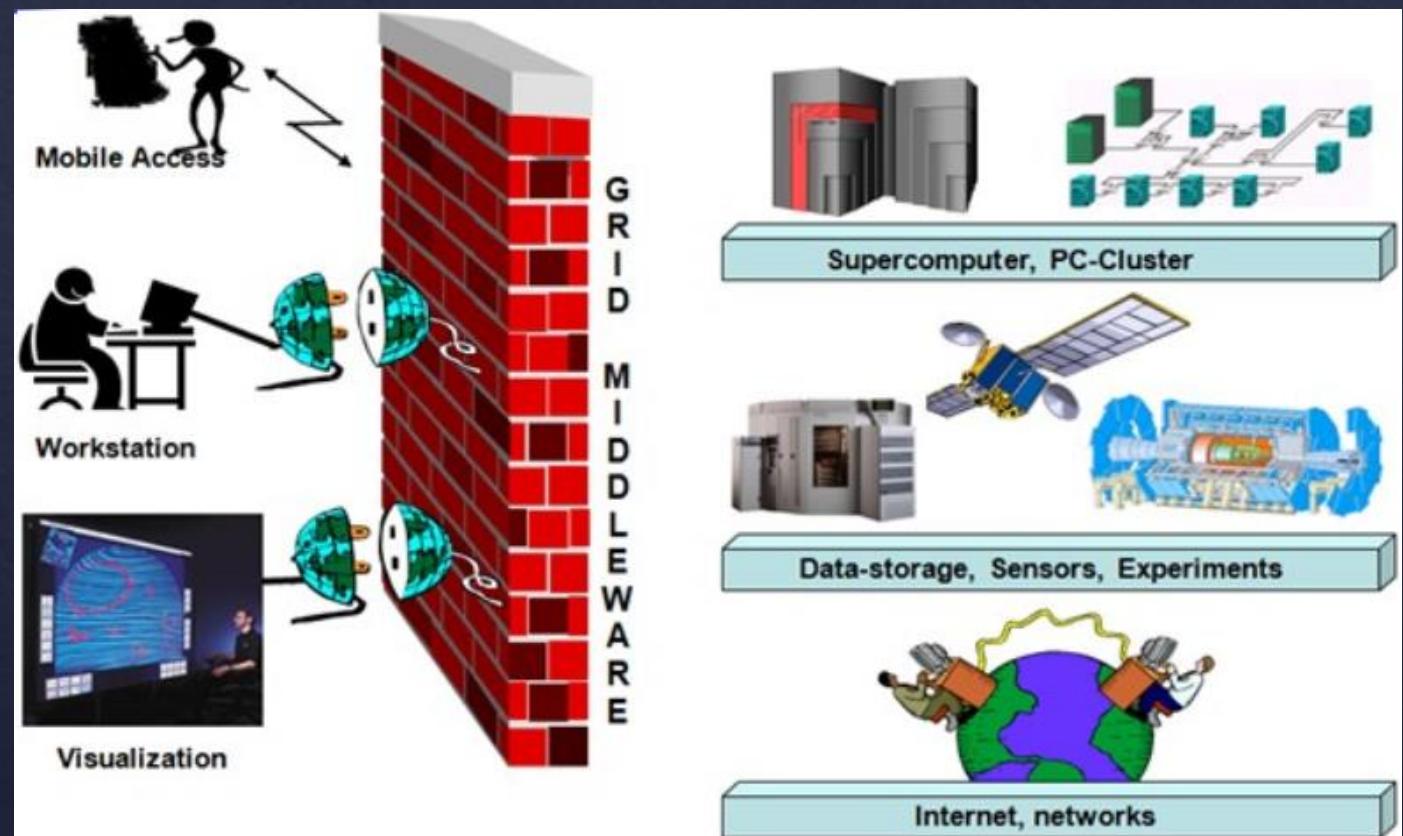
All Applications

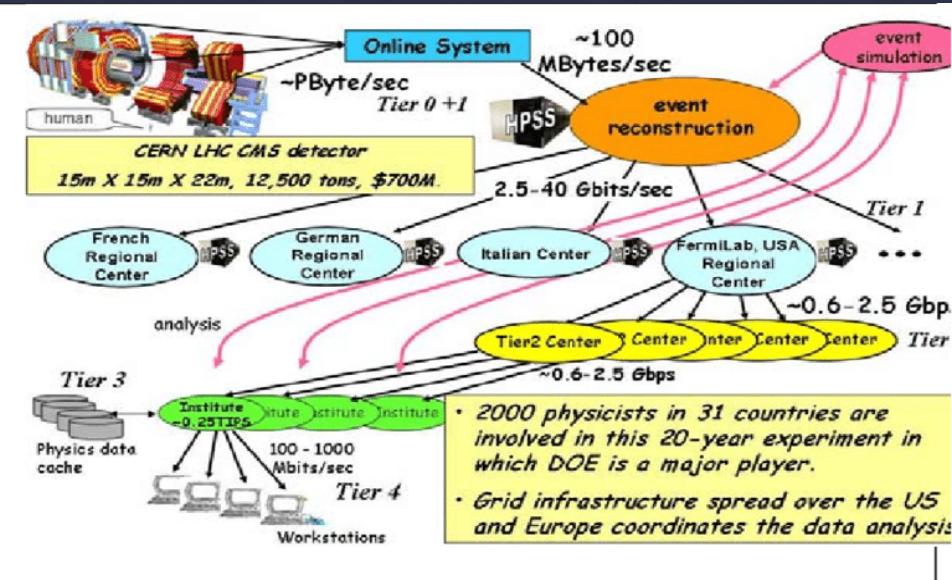


Three Main Issues in Grid Environment

- Heterogeneity
- Scalability
- Adaptability

All Applications





App



visPerf

Tools

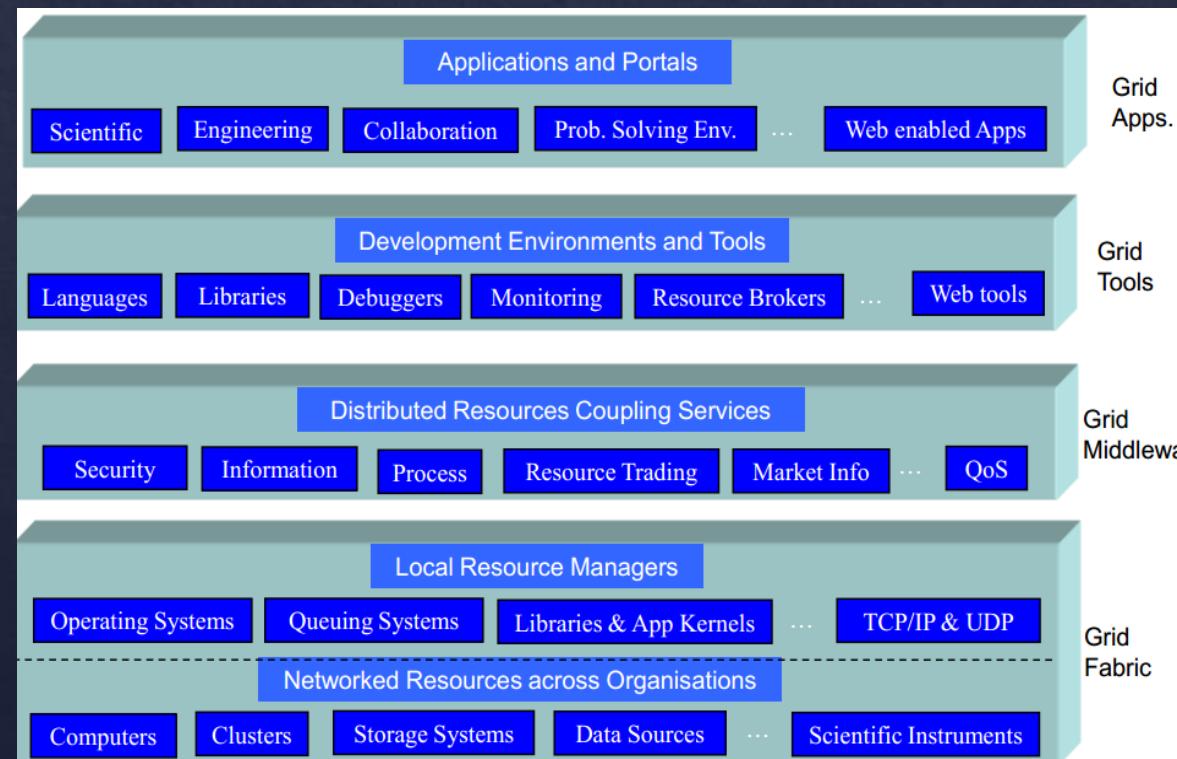


Monitoring



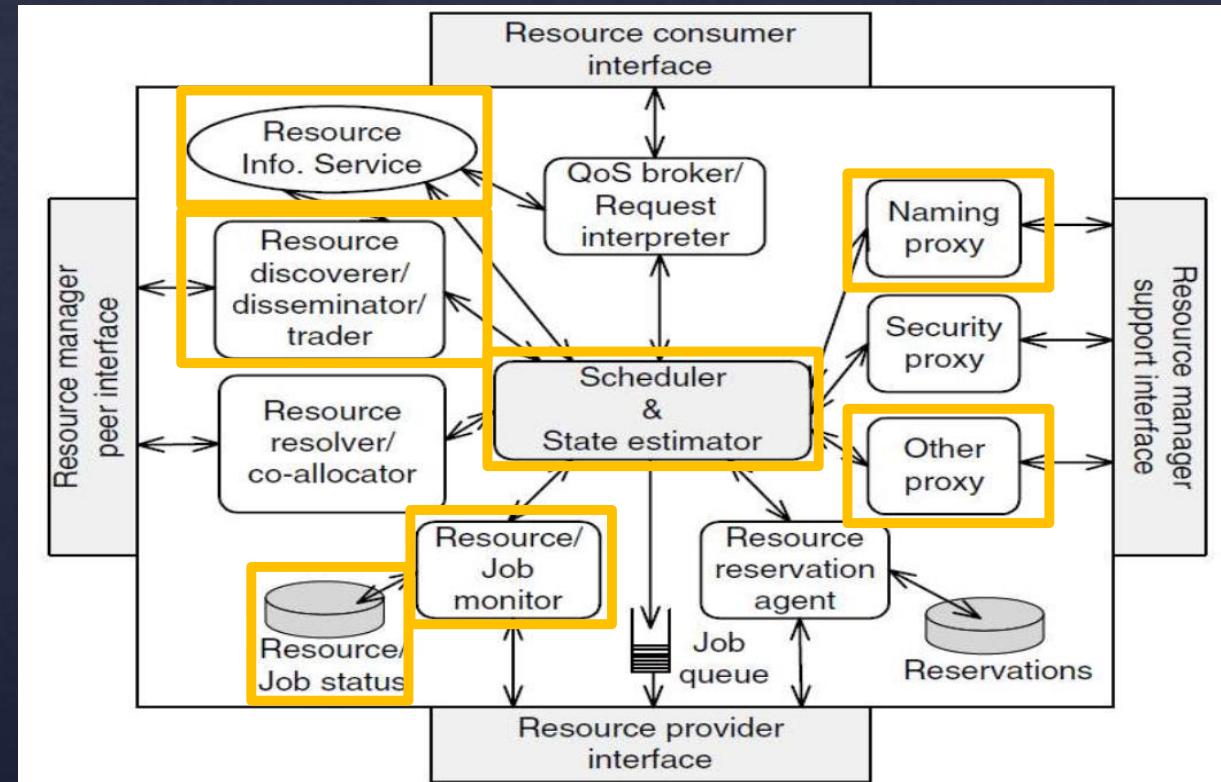
Fabric

Grid Software/Hardware Stack

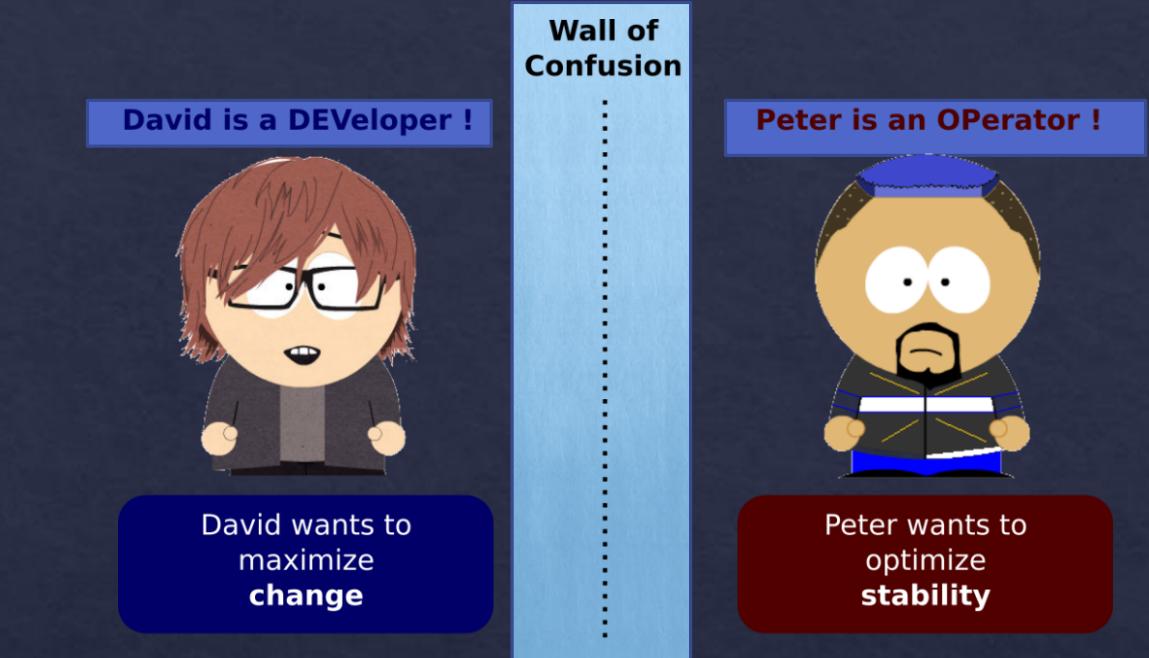
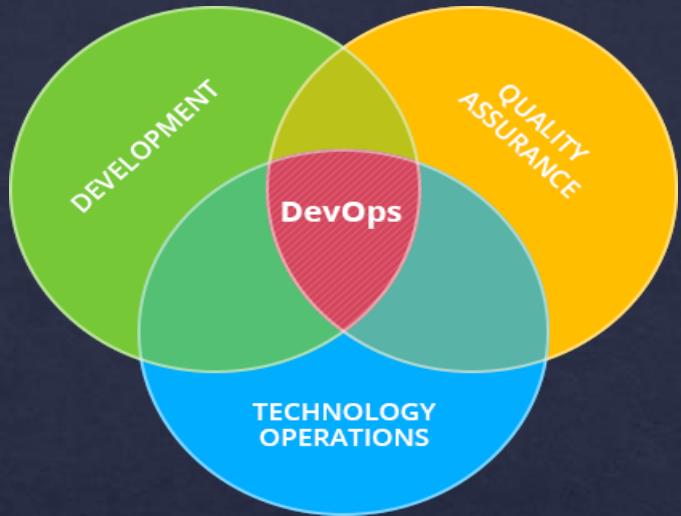


Resource Management System(RMS)

- Resource Dissemination and Discovery Protocols
 - Resource Dissemination Protocol
 - Information about resources
 - Discovery Protocol
 - Resource Discovery
- Resource Resolution and co-allocation Protocols
 - To
 - Schedule the job at the remote resources
 - To Simultaneously acquire multiple resources



Problem:



Solution: Continuous Integration/Continuous Development (CI/CD)



CI Combination of Tools:



Continuous Delivery

- Version Control Systems: Helps developers to maintain program source code onto the server.
 - Git
- Build Server: Retrieve source code periodically/automatically when the developer updates the code to VCS, and then trigger a new build
 - Jenkins
 - Travis CI
- Testing Automation Tools: Invokes the unit test program after the build succeeds, then notifies the result to the developer and QA.
 - Postman



Jenkins



Travis CI



POSTMAN

Configuration Management (CM):

- Helps to configure an OS Including:
 - The user
 - Group
 - System libraries
- Manages multiple servers keeping them consistent with the desired state or configuration



CM Combination of Tools:

- Ansible
- Puppet



Too Many Softwares

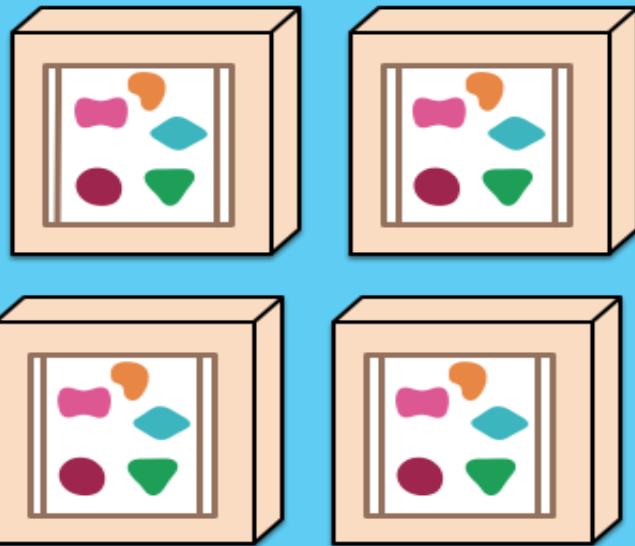
Monolithic Software Architecture:

- No Clear Definition
- But it Used to Have:
 - > than 50 modules or packages
 - > than 50 database tables
- May Takes
 - > than 30 minutes to build
- Cons
 - inflexible
 - unscalable
 - slow development
 - unreliable

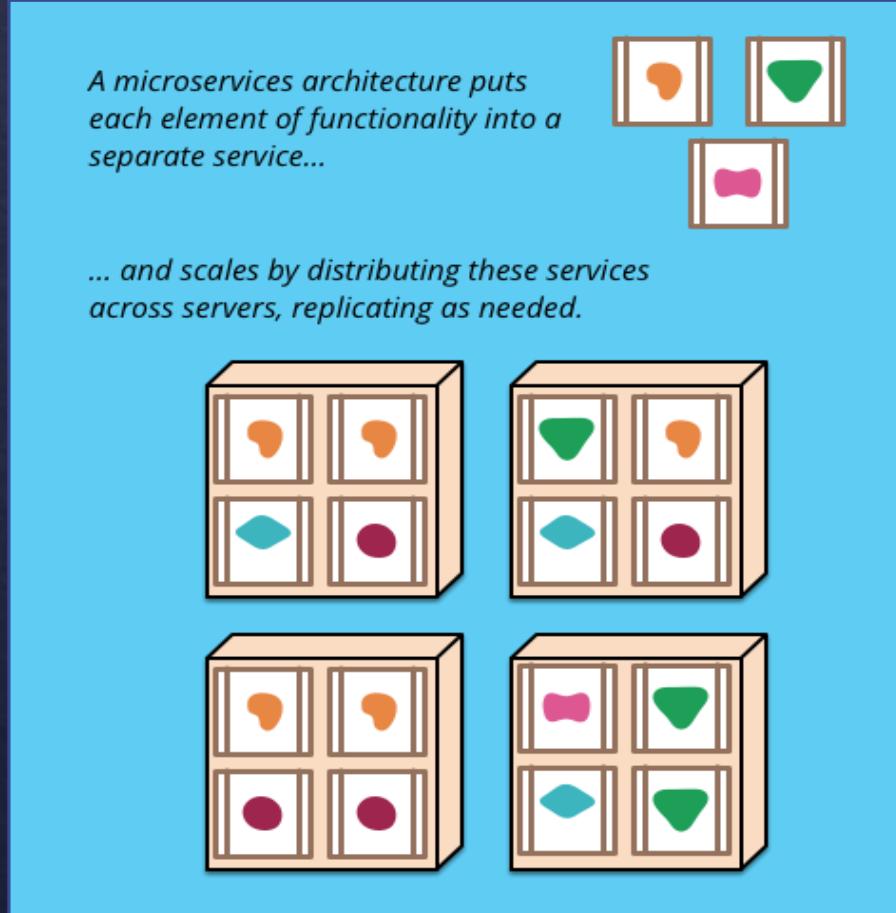
A monolithic application puts all its functionality into a single process...



... and scales by replicating the monolith on multiple servers



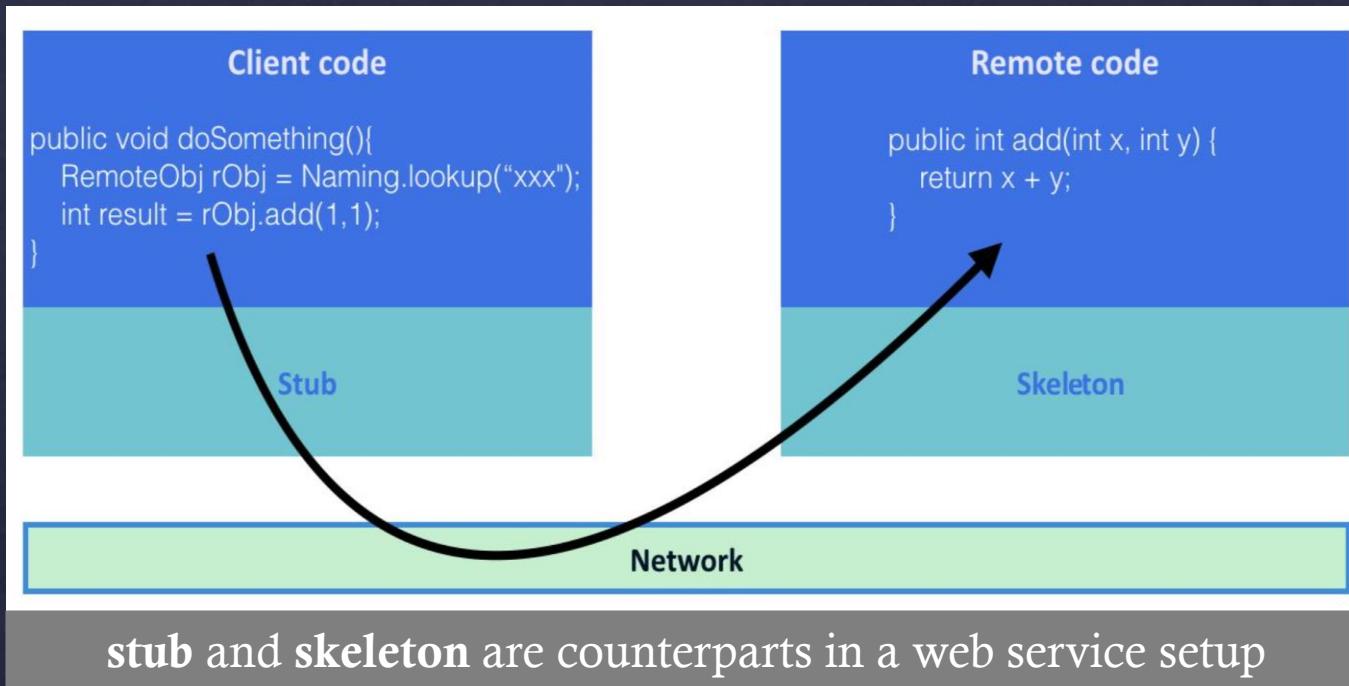
Microservice Software Architecture:



- Has been attempted back in the 1990s
- Structures an application as a collection of services that are:
 - Highly maintainable
 - Highly testable
 - Loosely coupled
 - Independently deployable

Remote Procedure Call:

- Use the module remotely (Good for Microservices)



- **Cons :**
 - Requires the **same programming language**
 - Less security consideration
- **Have the benefit :**
 - Divide application into **multiple processes/programs**
 - Programs can have **separate source code repositories**

RESTful design:

- Invented by [Roy Fielding](#)
- Defined in his PHD thesis in 2000
- Widely accepted by many developers
- De facto standard of distributed applications
- RESTful application allows any programming

language as it is HTTP-based

Roy Fielding



Roy Thomas Fielding (born 1965) is an American computer scientist, one of the principal authors of the HTTP specification, an authority on computer network architecture and co-founder of the Apache HTTP Server project.

Design - Projects

POST /design/projects Create a new item

GET /design/projects/{id} Find an item by ID

PUT /design/projects/{id} Update an item by ID

DELETE /design/projects/{id} Delete an item by ID

POST /design/projects/all Lists tests by ids

GET /design/projects/by-workspace/{workspaceId}/{type} List projects by workspace ID and type

Microservices:

- Should care about the following topics:

- **Stateless:**

- Doesn't store user session
 - Helps to scale out easier.

- **Versioning and compatibility:**

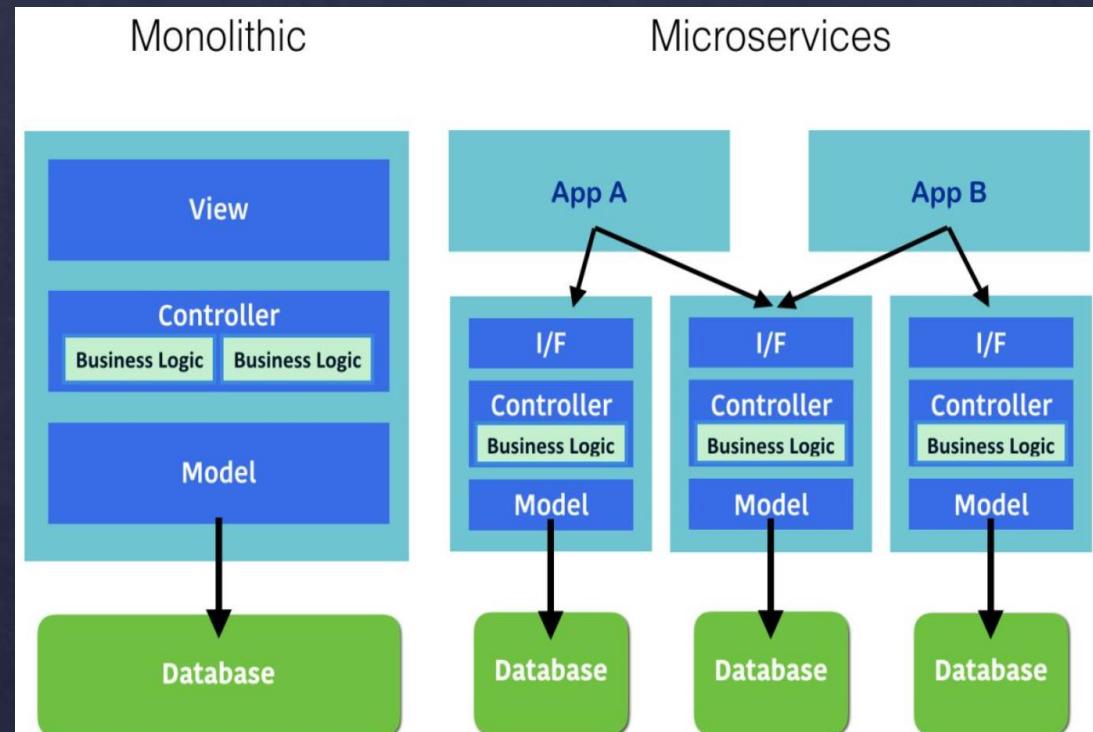
- May change and update the API
 - But should define a version
 - Should have backward compatibility

- **No shared datastore:**

- The microservice owns the datastore such as database.

- **Integrate CI/CD:**

- Should adopt CI/CD process to eliminate management effort.



Contents Cont.

5. Container Concepts and Fundamentals
6. Docker Hands On
7. Kubernetes Hands On



Development of Virtualization

- They say necessity is the mother of invention, and the history of computers is no exception
- Back in the '60s, computers were a rare commodity
 - To rent one cost well over a thousand dollars a month
 - which is why in the 1960s and through the 1970s, we saw the development of virtualization
- They made it possible to control all processing from a single location
 - Multiple computer terminals were connected to a single mainframe
- Disadvantages
 - If the user were to crash the central computer
 - The system would go down for everyone

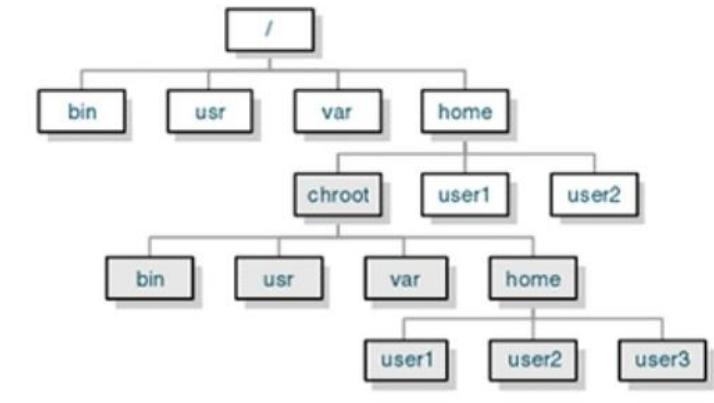


```
root@localhost.localdomain /var/log/yum.log (Fri Feb 21 21:57:45 2014) [0.130000]
File Edit View Search Terminal Help
http://www.tecmint.com
Sun Feb 16 03:46:04 2014 - main::HUP_handler: opening
o r midday ftp 0 * c
Wed Feb 19 01:04:25 2014 1 123.108.225.226 1771 /da
ta/httpd/htdocs/midday-beta/authrv2/mysql.php a _i
r midday ftp 0 * c
Wed Feb 19 01:05:12 2014 1 123.108.225.226 1770 /da
ta/httpd/htdocs/midday-beta/authrv2/mysql.php a _i
r midday ftp 0 * c
Wed Feb 19 01:05:50 2014 1 123.108.225.226 158 /dat
a/httpd/htdocs/midday-beta/authrv2/orc_config.php a
i r midday ftp 0 * c
Wed Feb 19 01:06:44 2014 1 123.108.225.226 13597 /da
ta/httpd/htdocs/midday-beta/authrv2/search.php a
i r midday ftp 0 * c
Wed Feb 19 01:07:20 2014 1 123.108.225.226 13601 /da
ta/httpd/htdocs/midday-beta/authrv2/search.php a
i r midday ftp 0 * c
Wed Feb 19 01:08:12 2014 1 123.108.225.226 13347 /da
ta/httpd/htdocs/midday-beta/authrv2/search.php a
i r midday ftp 0 * c
Wed Feb 19 01:12:04 2014 1 123.108.225.226 1682 /ho
me/httpd/htdocs/cms-beta/oracleapi/mysql.php a _o
r midday ftp 0 * c
Wed Feb 19 01:12:23 2014 1 123.108.225.226 1618 /ho
me/httpd/htdocs/cms-beta/oracleapi/mysql.php a _i
r midday ftp 0 * c
Wed Feb 19 01:12:57 2014 1 123.108.225.226 1582 /ho
me/httpd/htdocs/cms-beta/oracleapi/mysql.php a _i
r midday ftp 0 * c
Feb 21 20:22:47 Installed: 2:vim-common-7.2.411-1.8.e
02] /var/log/ajenti.log *Press F1/<CTRL>+<h> for help*
.org
2013-10-13 21:53:57,268 INFO connectionpool.new
conn(): Starting new HTTP connection (1): meta.ajenti
.org
2013-10-14 09:53:59,456 INFO connectionpool.new
conn(): Starting new HTTP connection (1): meta.ajenti
.org
2013-10-14 21:54:01,559 INFO connectionpool.new
conn(): Starting new HTTP connection (1): meta.ajenti
.org
Feb 21 20:22:02 Erupted: multitail
Feb 21 21:22:25 Installed: multitail-5.2.13-1.el6.rf.
1686
Feb 21 20:23:13 Installed: 2:vim-X11-7.2.411-1.8.el6.
1686
Feb 21 20:23:14 Installed: parcellite-0.9-1.el6.rf.16
86
00] /var/log/xferlog *Press F1/<CTRL>+<h> for help*
00] /var/log/yum.log 954 - 2014/02/21 21:57:46
```

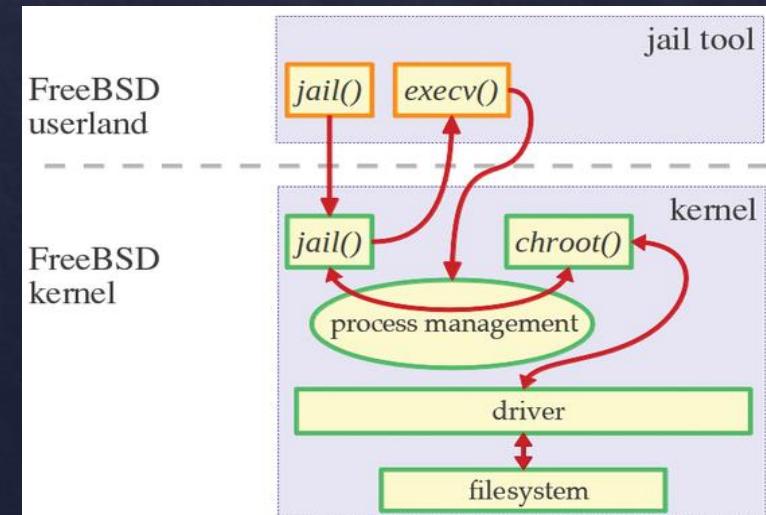
Development of Virtualization

- We took another step towards creating shared, yet isolated, in 1979
 - Development of the chroot (change root) command
 - Isolates system processes into their own segregated filesystems

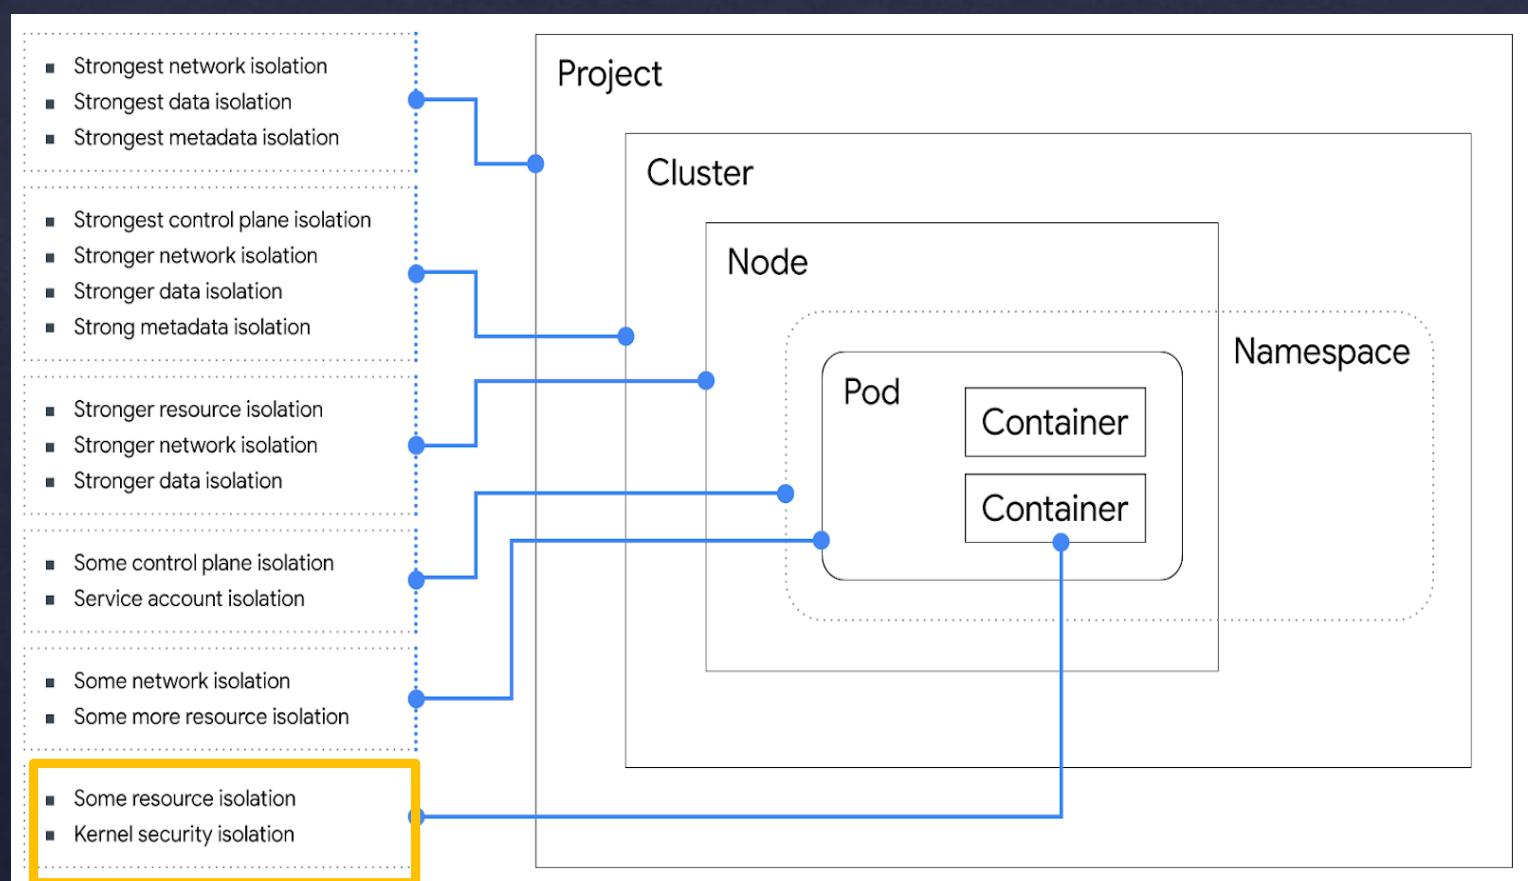
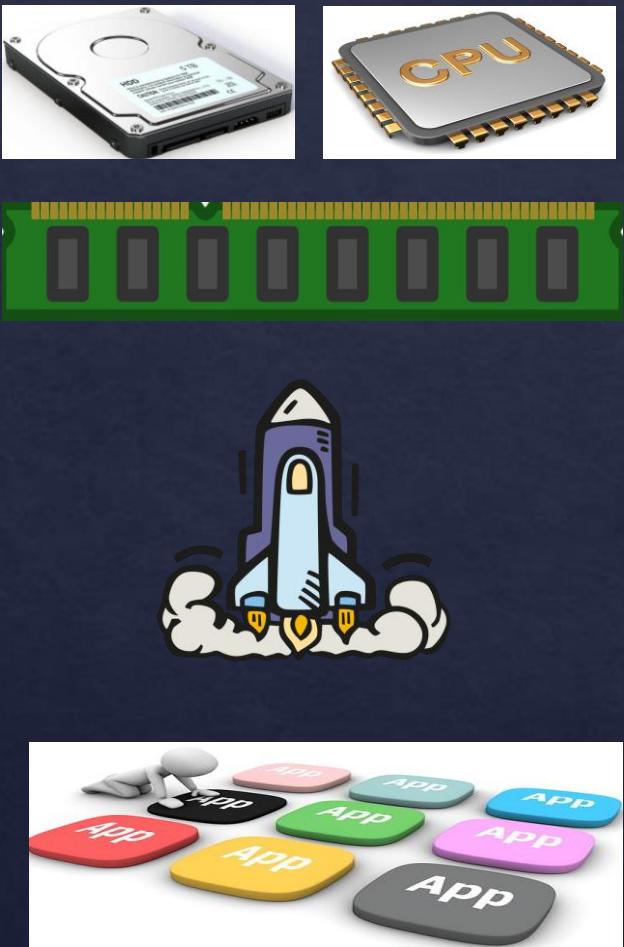
chroot



- Skip forward a bit in time to the **1990s**
 - Bill Cheswick, His solution was to make modifications to a chrooted environment
 - Result of his studies → Linux jail command.
 - Process sandboxing features:
 - ✓ File systems
 - ✓ Users
 - ✓ Networks
 - ✓ etc.



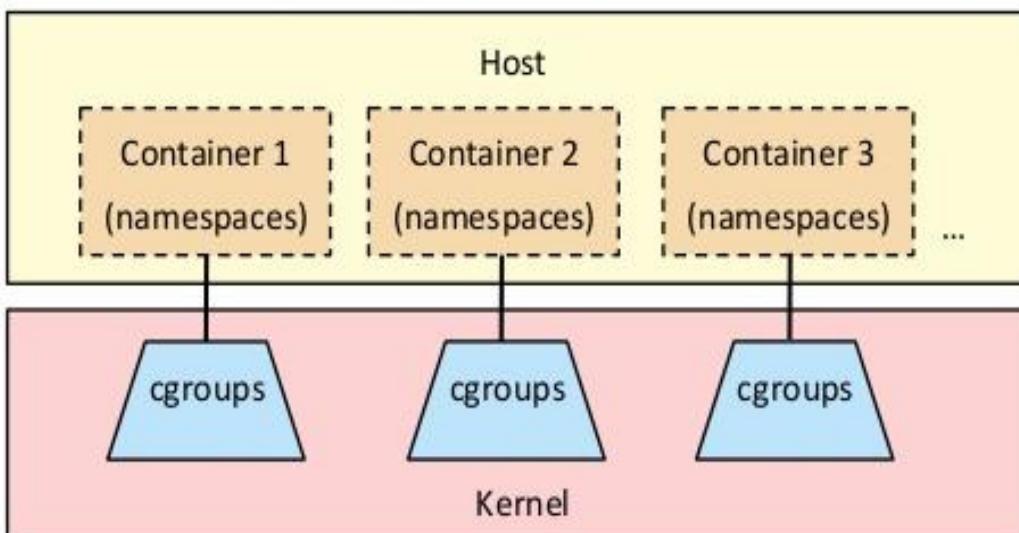
The key feature of container is isolation



Linux Namespaces

- Provide logical partitions of system resources
 - mounting point (mnt)
 - process ID (PID)
 - network (net)

Container = combination of namespaces & cgroups



```
$ ps axf
 PID TTY      STAT   TIME  COMMAND
  2 ?        S      0:00  [kthreadd]
  3 ?        S      0:42  \_ [ksoftirqd/0]
  5 ?        S<     0:00  \_ [kworker/0:0H]
  7 ?        S      8:14  \_ [rcu_sched]
  8 ?        S      0:00  \_ [rcu_bh]    1
```

```
$ sudo unshare --fork --pid --mount-proc=/proc /bin/sh
$ ps axf
 PID TTY      STAT   TIME  COMMAND
  1 pts/0     S      0:00  /bin/sh
  2 pts/0     R+     0:00  ps axf
```

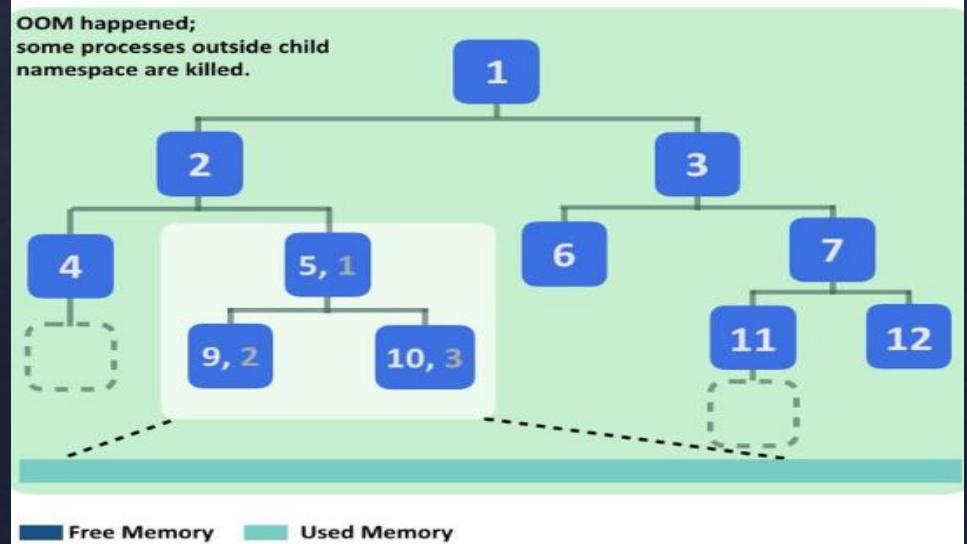
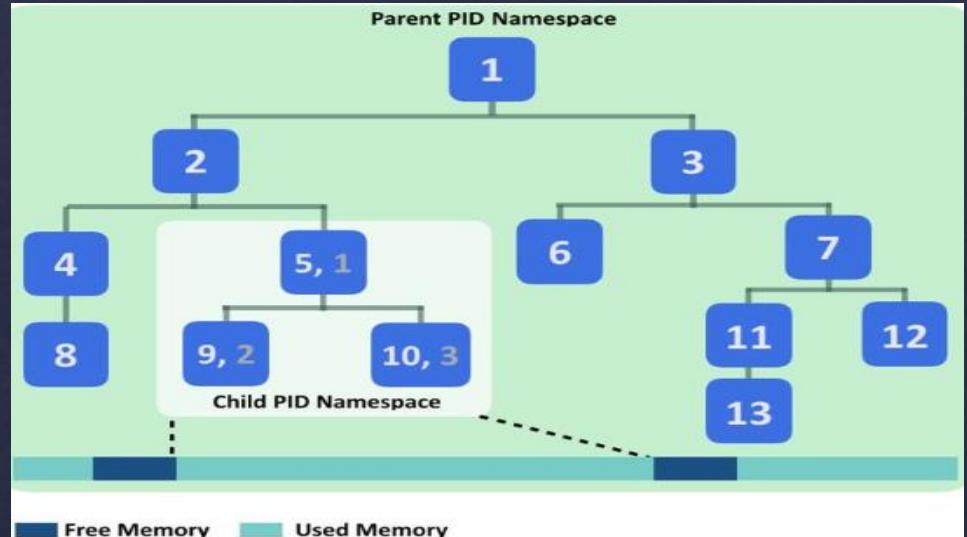
2

```
$ ps axf // from another terminal
 PID TTY      COMMAND
 ...
25744 pts/0 \_ unshare --fork --pid --mount-proc=/proc
/bin/sh
25745 pts/0     \_ /bin/sh
3305 ?        /sbin/rpcbind -f -w
6894 ?        /usr/sbin/ntpd -p /var/run/ntp.pid -g -u
```

3

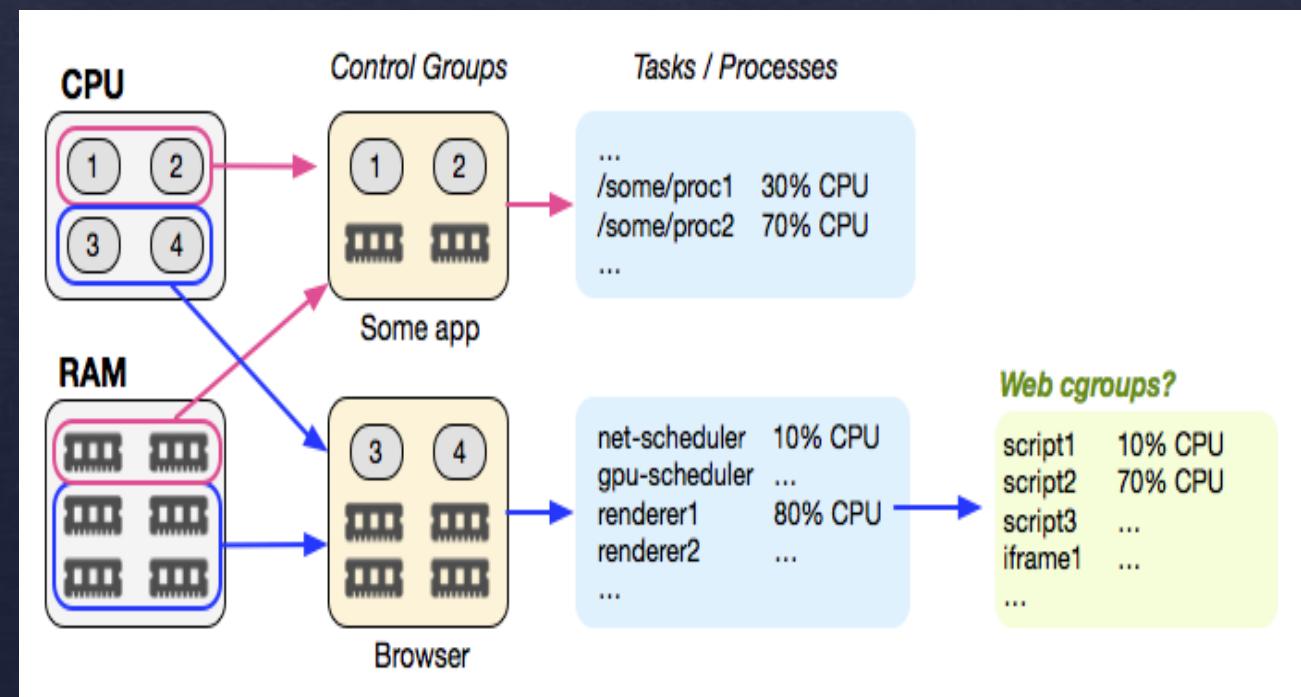
Out of Memory

The following diagram illustrates the PID namespaces and how an out-of-memory (OOM) event can affect other processes outside a child namespace

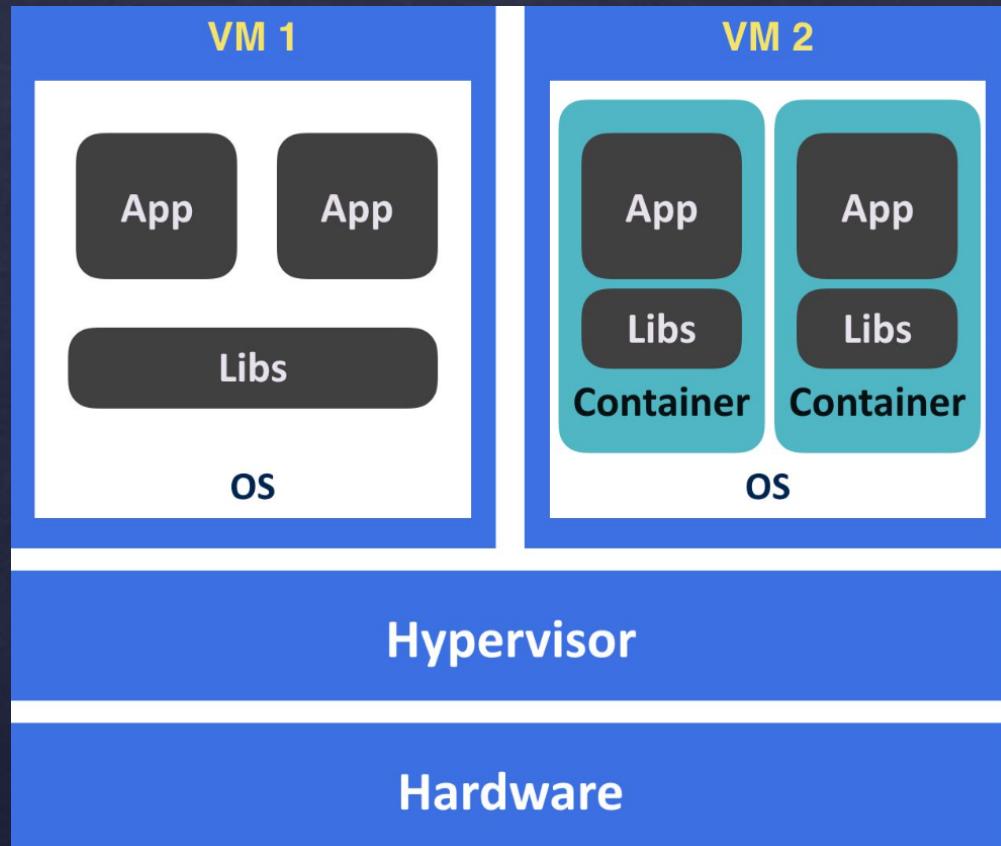


Linux Control Groups(cgroup)

- Uses in order to
 - Limit resource usage by
 - Setting constraint on different kinds of system resources



These Days of Virtualization

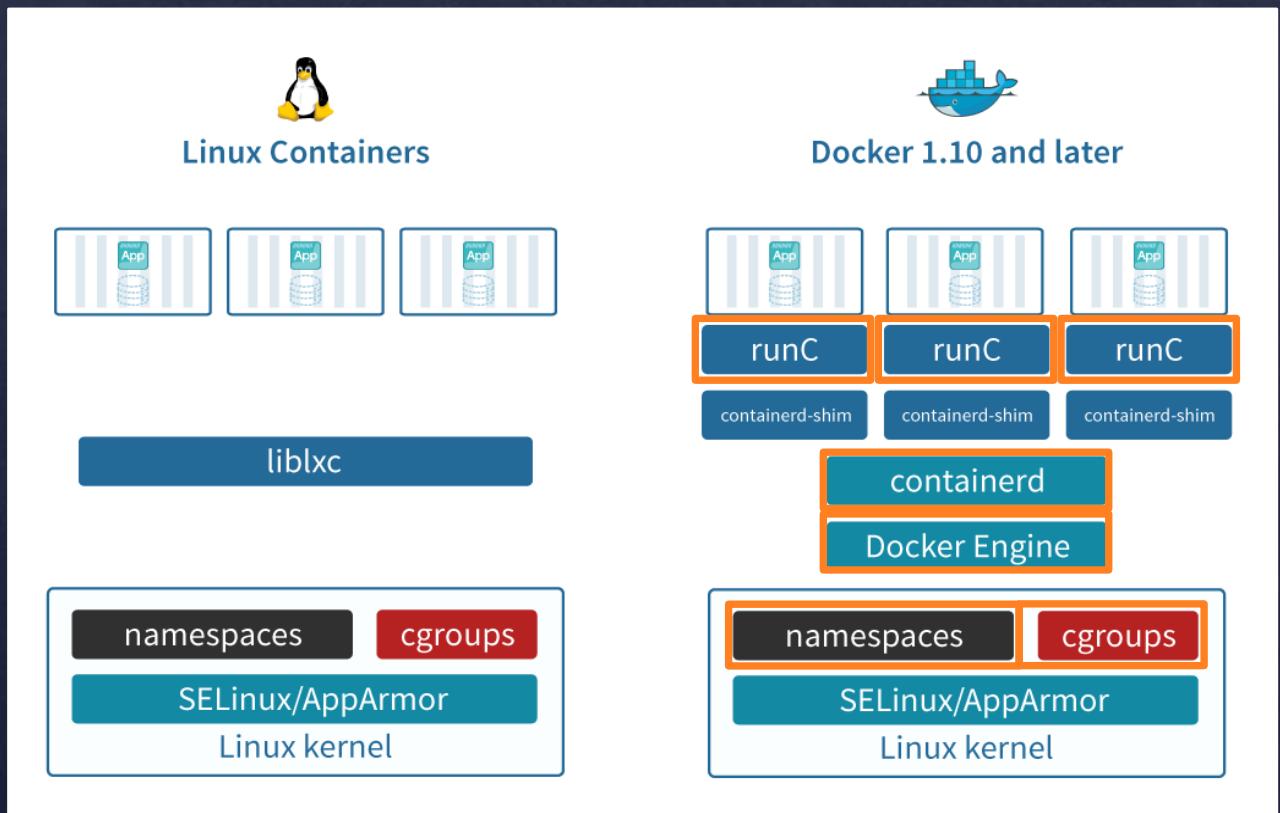


Terms:

- Virtualization
- Containerization

Linux Container

- Two most important building blocks
 - namespaces
 - cgroups (control groups)



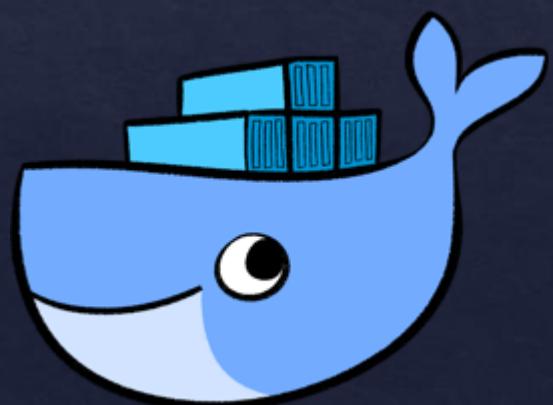
Contents Cont.

5. Container Concepts and Fundamentals
6. Docker Hands On
7. Kubernetes Hands On



Docker Fundamental

- An Open-Source GO framework
- Solomon Hykes starts Docker as an internal project within dotCloud in 2013
- Whale is the one who won the contest



What's in a name?

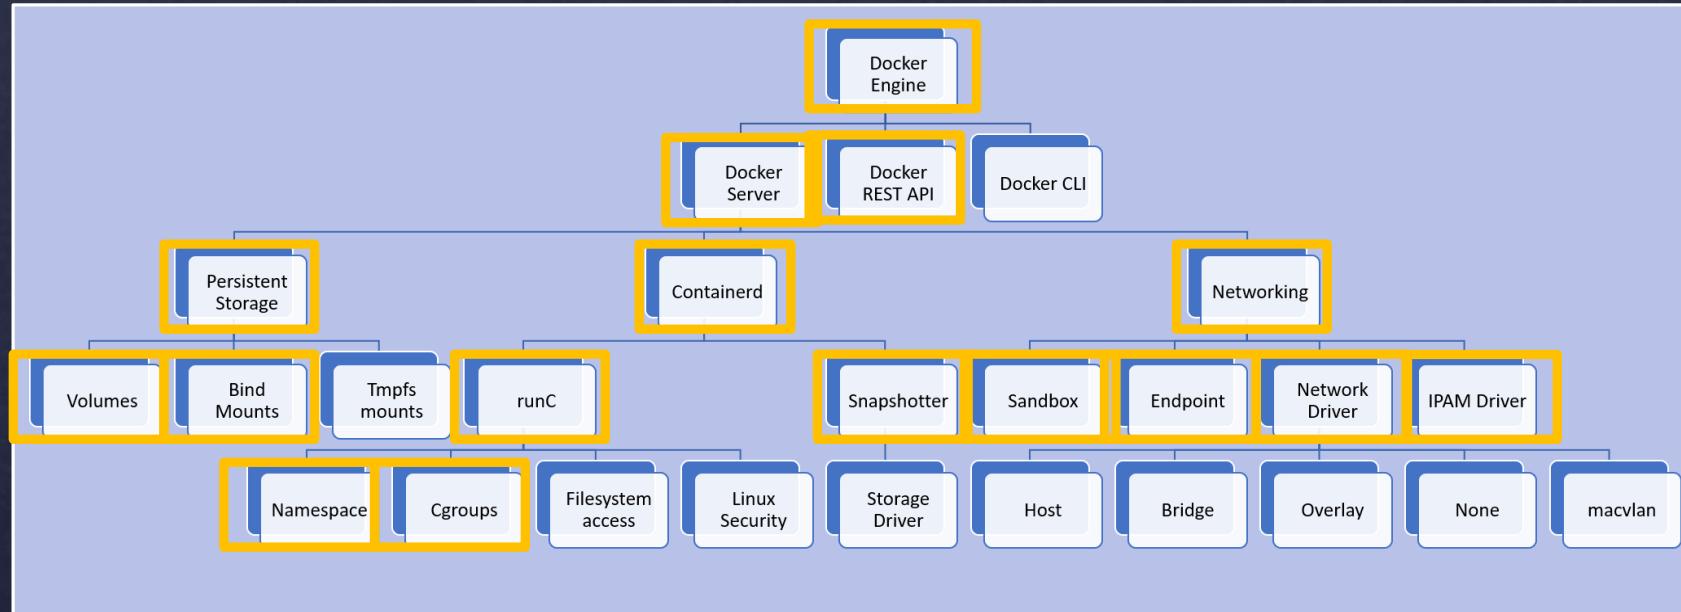


Why GO



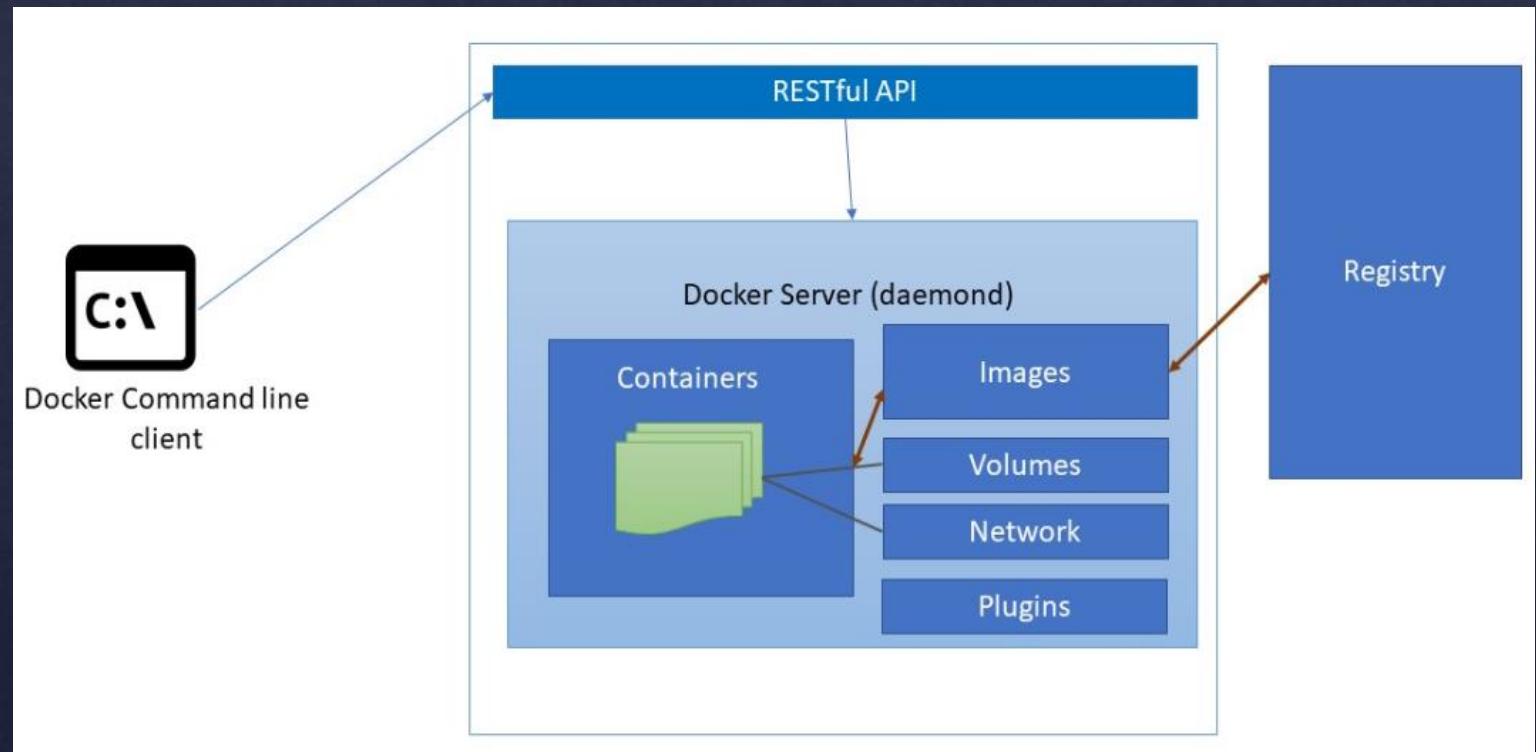
- Simplicity is Complex
- Medium By Example

Overall Hierarchy of Docker



Docker Engine

- The heart of Docker is the docker engine
- Consists of:
 1. Docker Server
 2. RESTful API
 3. Docker CLI



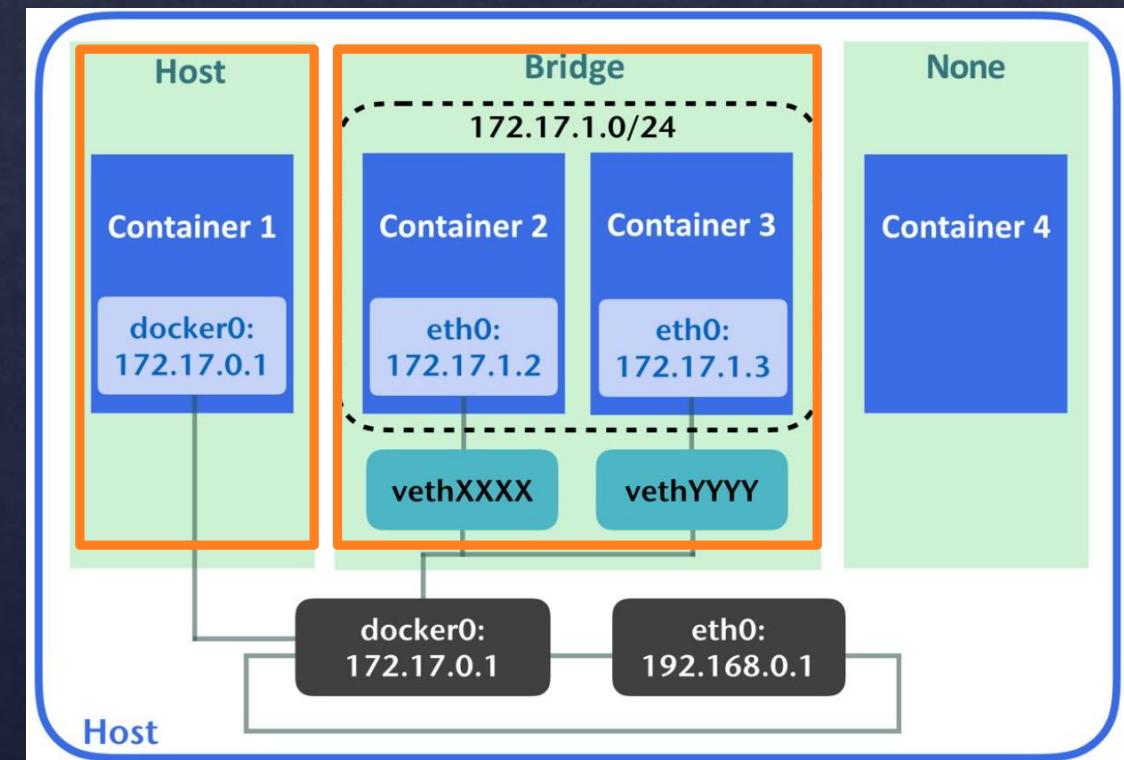
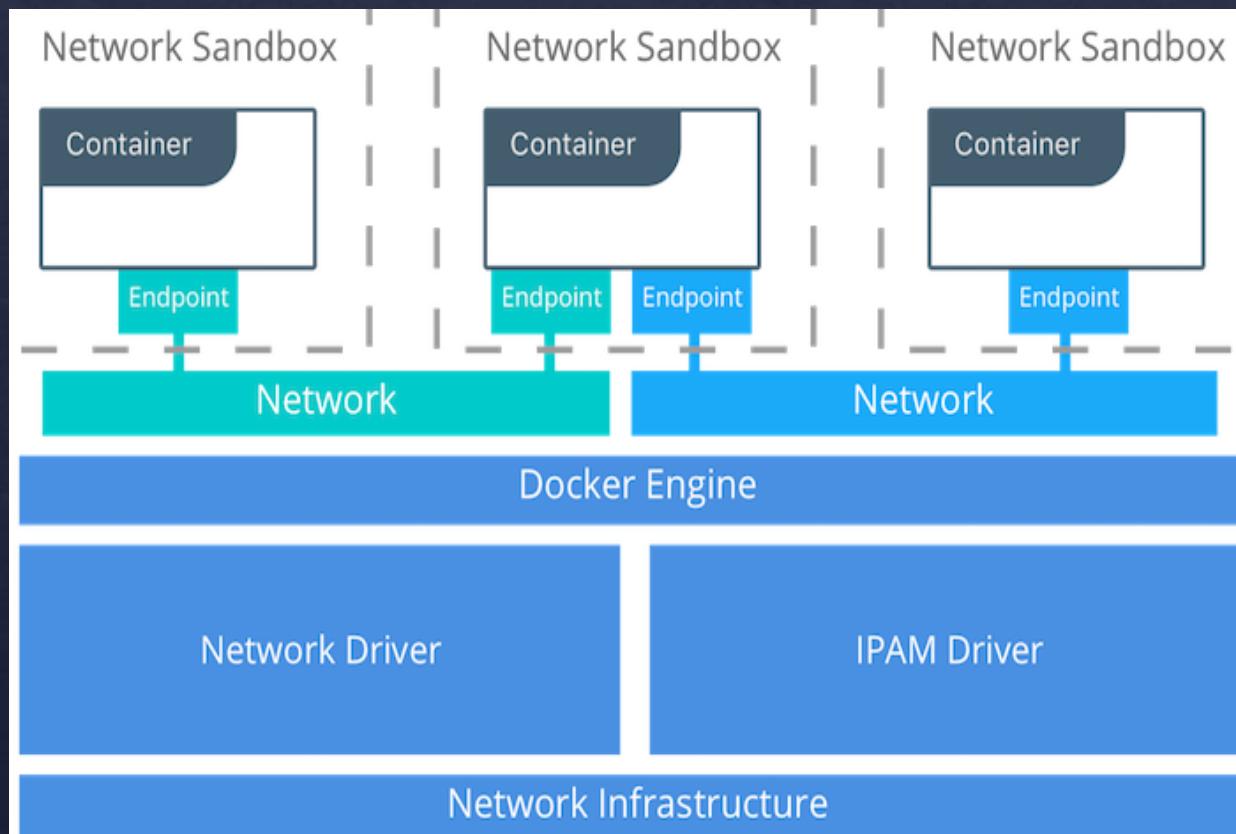
Persistent Storage

- It is used for making data persistent outside the container
- Consists of:
 - Volumes
 - Maintaining volumes
 - ✓ They are managed using the API
 - ✓ They can be shared amongst multiple containers
 - ✓ They can work on both linux, windows or even remote hosts or cloud providers
 - Bind Mounts
 - You can mount a file or directory from the host machine to the docker container using bind mounts
 - tmpfs Mounts
 - This can be used to store temporary data
 - Data is kept on the temporary storage area of the host machine.

Networking

- Consists of:
 - **Sandbox:** It manages the container...
 - Routing tables
 - Interfaces
 - DNS
 - **Endpoint**
 - Endpoints Join a Sandbox to a network.
 - **Network Driver:** used by the docker engine to connect to the actual network infrastructure
 - Remote: Container uses Host networking
 - Native: **Host, Bridge, Overlay, None ...**
 - **IPAM Driver:** Manages IP addresses for Docker

Networking



Containerd: The Heart of Container System

- Docker Engine uses it to create and manage Containers
- It **abstracts calls to system or OS specific functionality** to run containers on windows, solaris and other Oss
- The scope of containerd includes:
 - Create, start, stop, pause, delete a container, ...
 - Functionality for copy on write file systems for containers (**Copy on write**)
 - Build, push and pull images and management of images
 - Persisting container logs

Containerd: The Heart of Container System CONT.

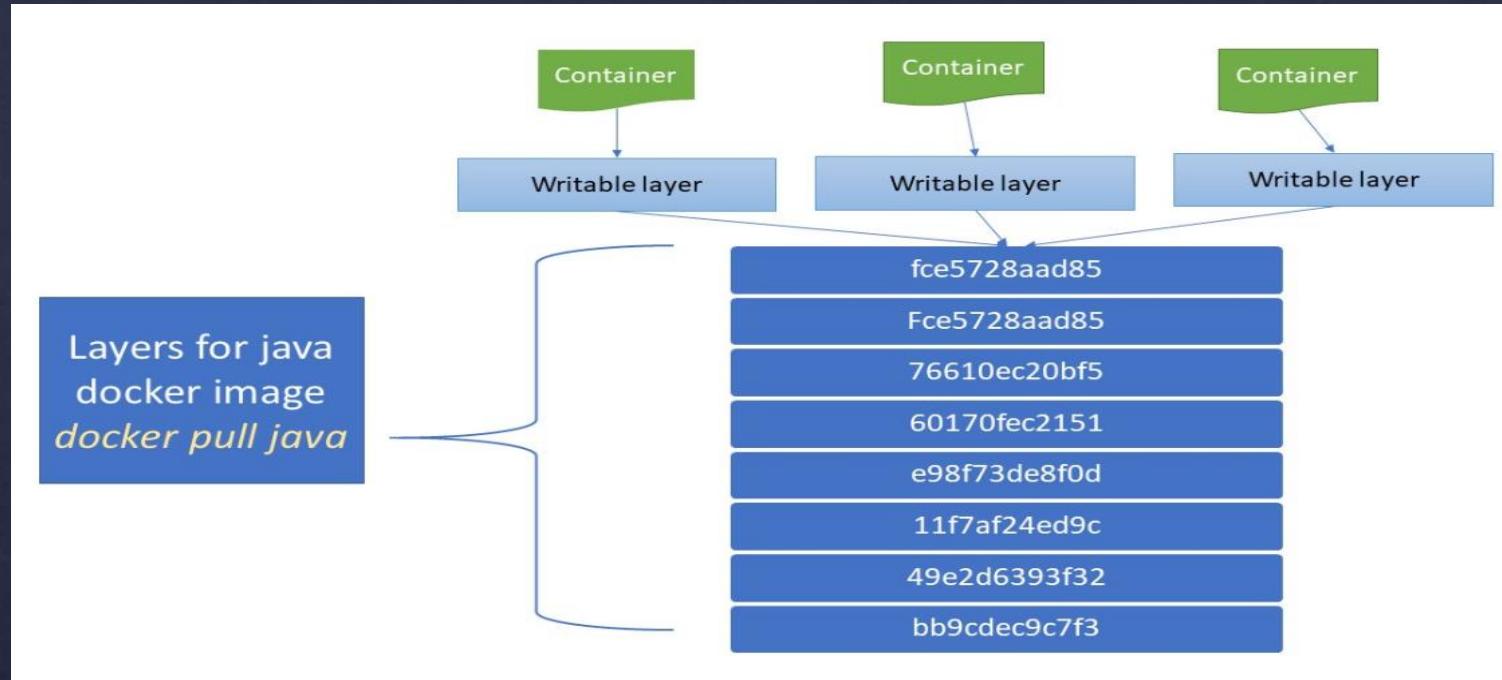
- Docker Engine uses it to create and manage Containers

- Consists of:

- Container runtime – runC
 - runC is a cli tool that follows the open container initiative(OCI)
 - runC provides GO implementation to create containers using
 - ✓ Namespace
 - ✓ Cgroups
 - ✓ Filesystem access controls
 - ✓ Linux security capabilities
- Snapshotter
 - A snapshot is a filesystem state
 - Docker containers use a system known as layers
 - Layers allow making modifications to a file system and storing them as a changeset on the top of a base layer
 - A layer is the diff between the snapshots

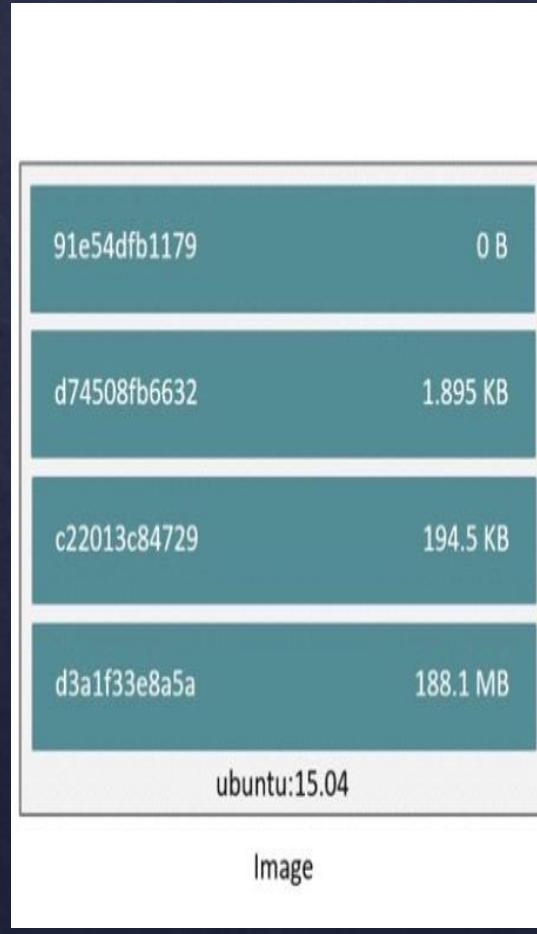
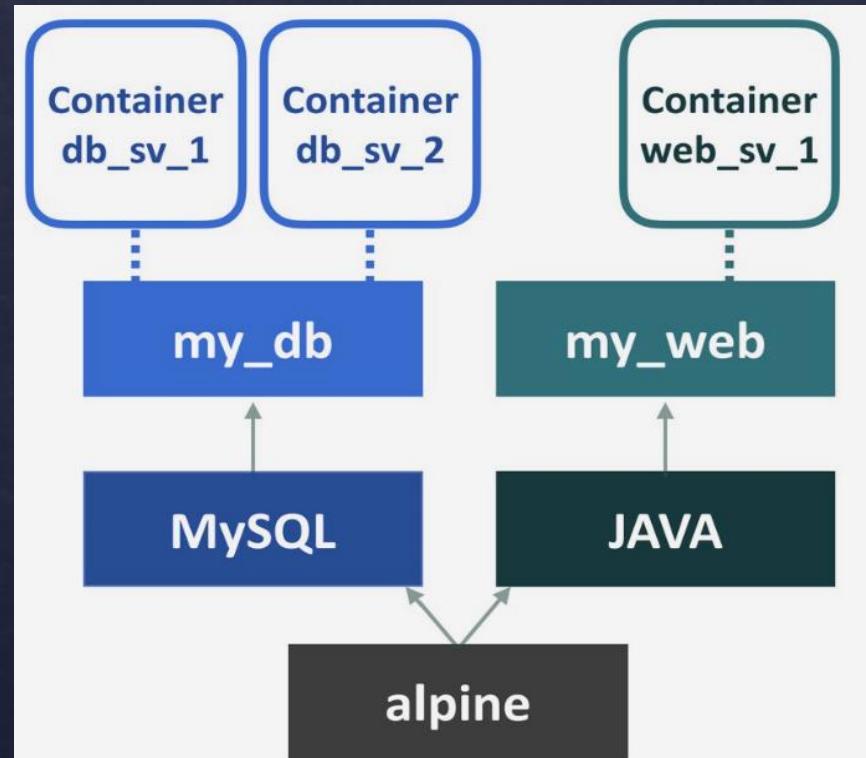


Snapshooter



- When a container is created
 - Adds a **writable layer(Copy on Write)** on top of all the layers
 - All changes are written to this writable layer
 - This writable layer is what differentiates a container from an image.

The Layers



Container
(based on ubuntu:15.04 image)

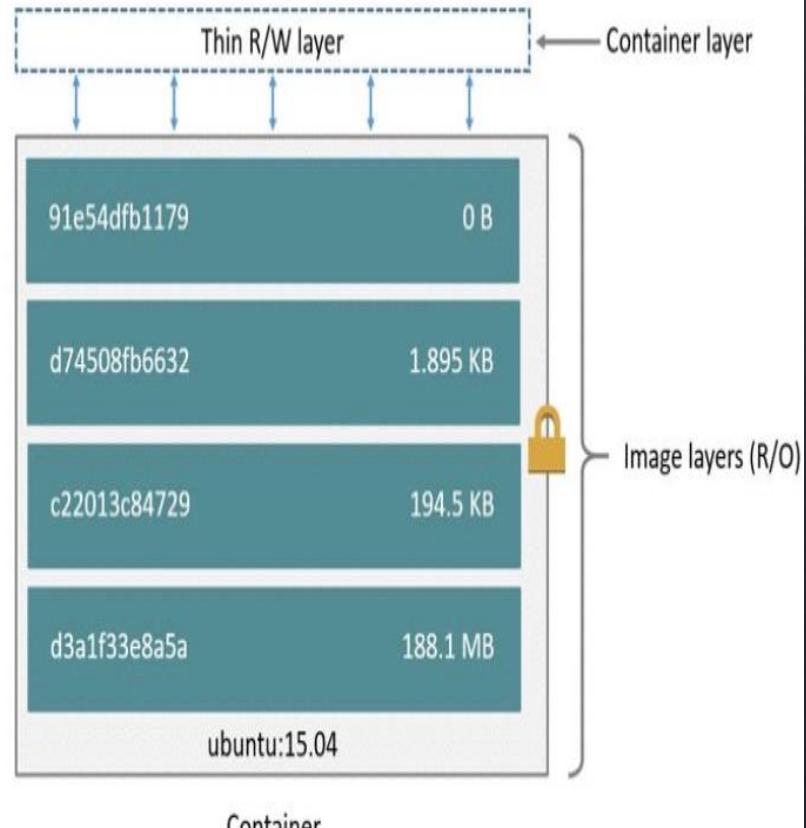
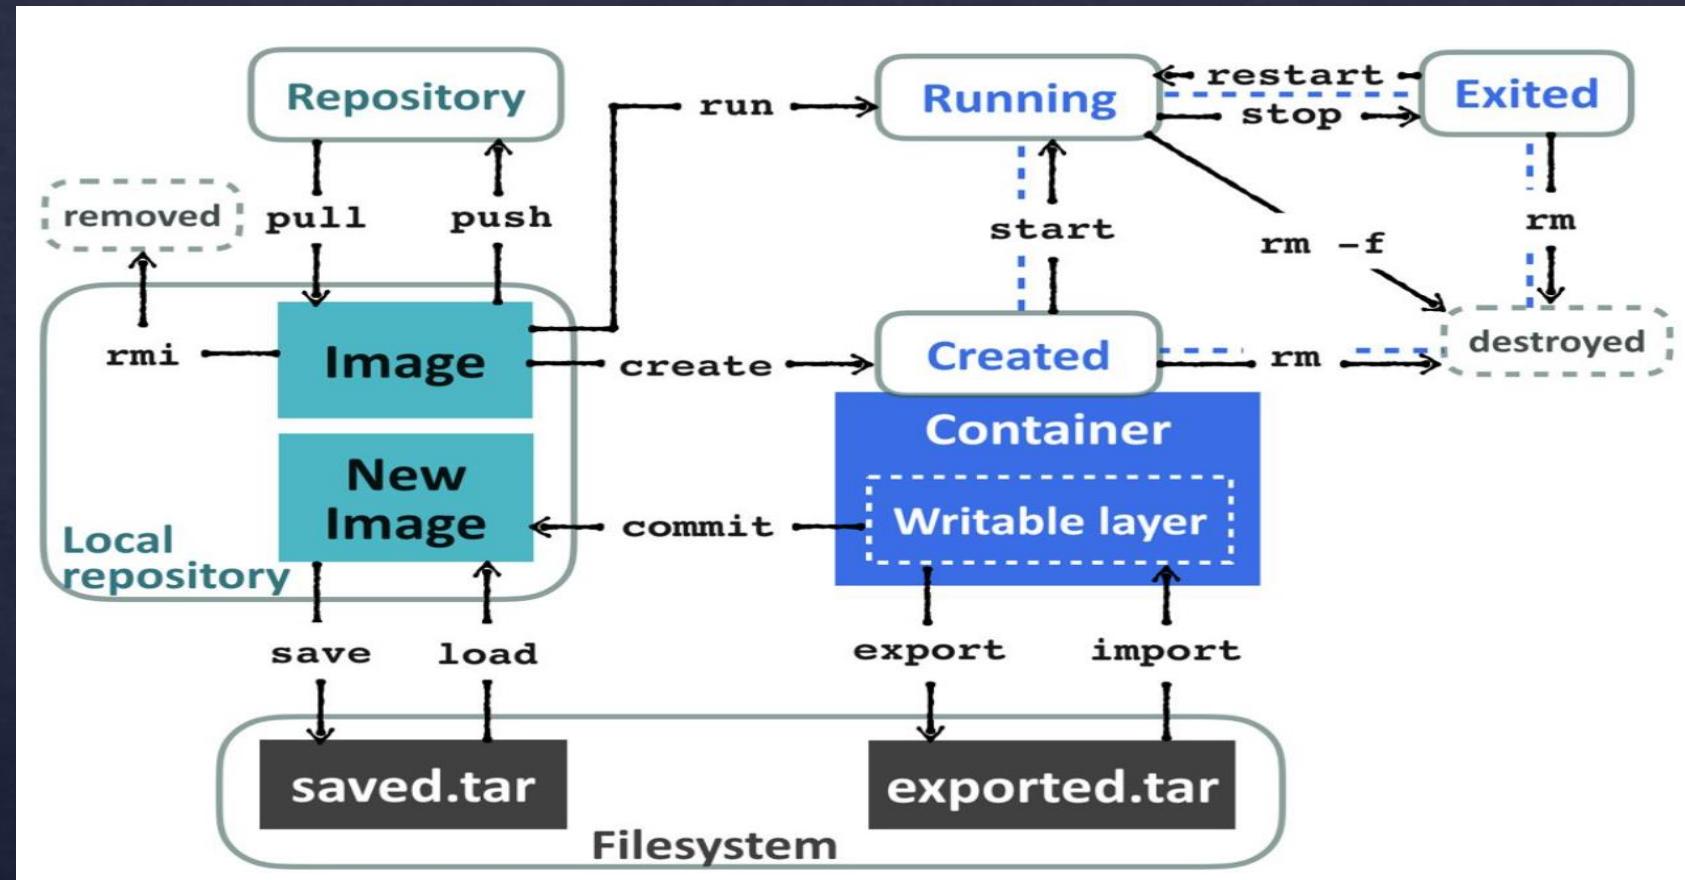


Image Vs Container

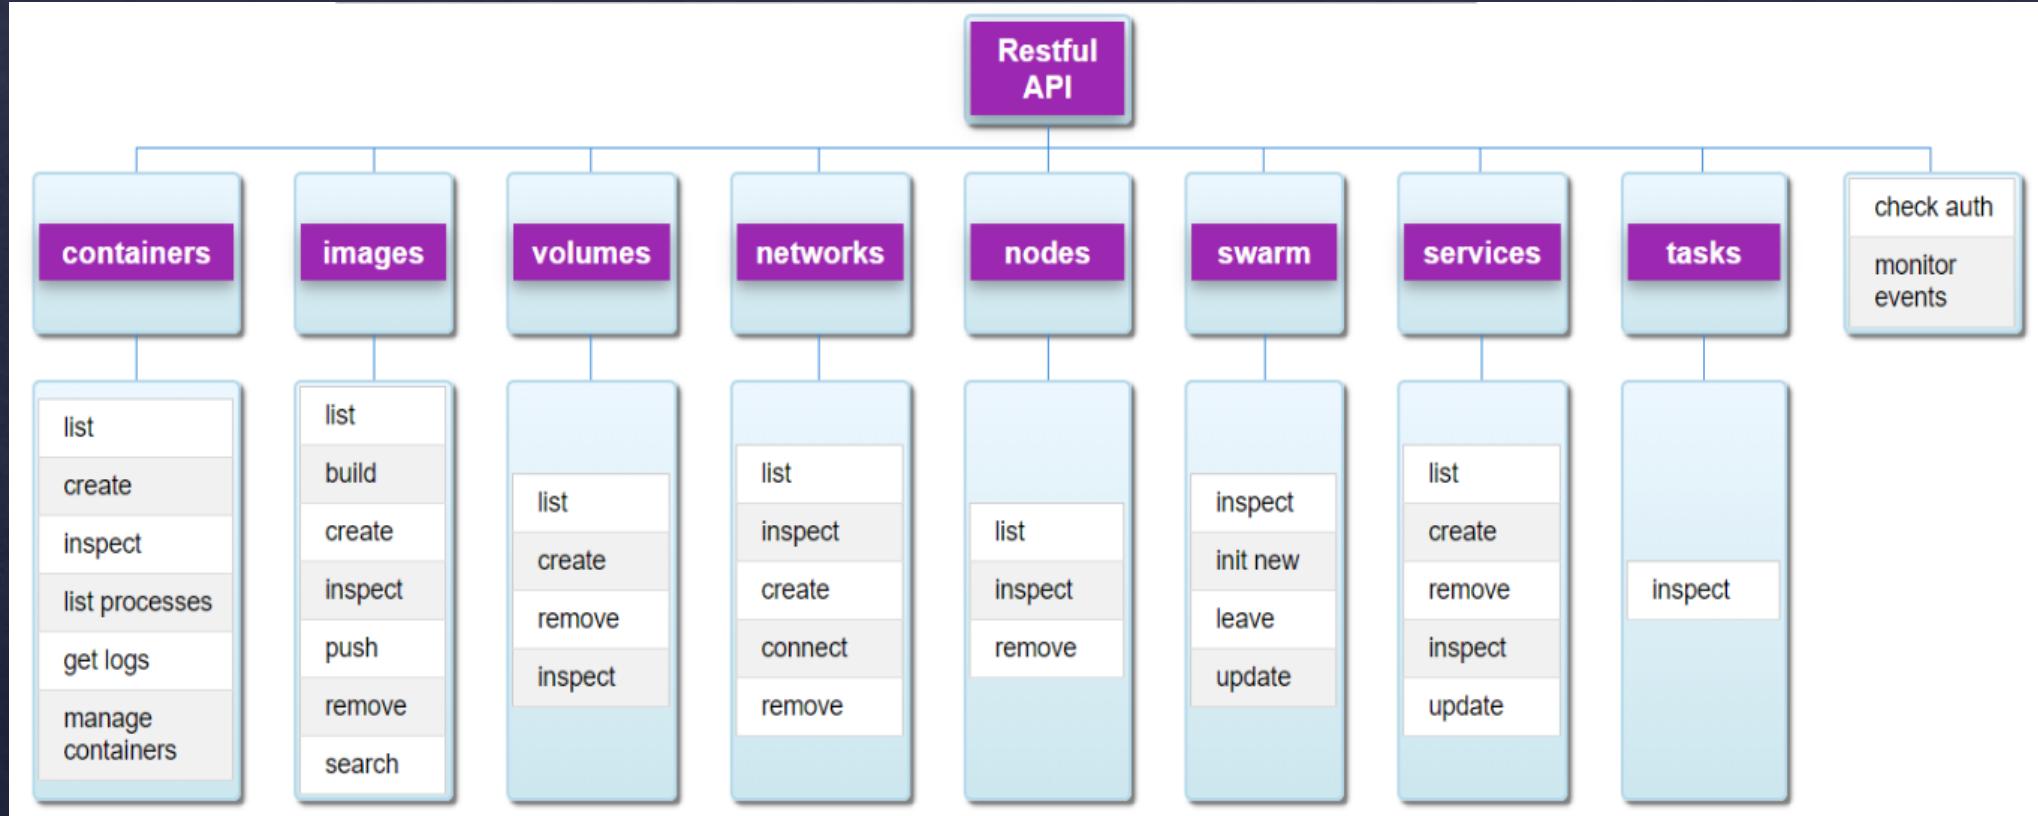
The following diagram depicts the relationship of states between container and images. The captions on the arrows are the corresponding sub-commands of Docker



The Snapshotter Commands

- `docker commit [CONTAINER]`: Commits the changes of the container layer into a new image
- `docker save --output [filename] IMAGE1 IMAGE2 ...`: Saves one or more images to a TAR archive
- `docker load -i [filename]`: Loads a tarball image into the local repository
- `docker export --output [filename] [CONTAINER]`: Exports a container's filesystem as a TAR archive
- `docker import --output [filename] IMAGE1 IMAGE2`: Imports a filesystem tarball

Docker RESTful API



❖ Hands On Docker/Overall Hierarchy of Docker/Docker Engine/**RESTful API/1**

Glossary

- **Docker Client:** the binary interacting with the Docker Engine.
- **Docker Image:** A filesystem (read-only template) used to create a Container (think “the binary”)
- **Docker Container:** a running image providing a service (think “the process”)
- **Host:** The computer running the Docker Engine
- **Docker Registry:** A private or public (Docker Hub) collection of Docker Images
- **Docker Hub:** A cloud-based repository in which Docker users and partners create, test, store and distribute Docker Images.
- **Docker Machine:** Provision hosts and install Docker on them
- **Docker Compose:** Create and manage multi-container architectures
- ***Docker Swarm:** Orchestrating tool to provision and schedule containers

Contents Cont.

5. Container Concepts and Fundamentals
6. Docker Hands On
7. Kubernetes Hands On



Kubernetes Fundamental

- An Open-Source **GO** framework
- Kubernetes (“**koo-burr-NET-eez**”) is
 - Conventional pronunciation of a **Greek** word
 - ✓ “κυβερνήτης” == **Helmsman/Pilot**



Getting Started with Kubernetes

- Questions:
 1. Is there a way to build services across multiple machines without dealing with cumbersome network and storage settings ?
 2. Is there any other easy way to manage and roll out our microservices by different service cycle ?

Understanding Kubernetes

- Kubernetes is a **platform** for managing application containers across multiple hosts.
- Same as the nature of containers: **It's designed to run anywhere**
 - Clusters
 - Grids
 - Cloud (public, private, hybrid)

Understanding Kubernetes

- Kubernetes considers most of the operational needs for application containers
 - Auto-scaling
 - Container deployment
 - Persistent storage
 - Container health monitoring
 - Compute resource management
 - High availability

Understanding Kubernetes

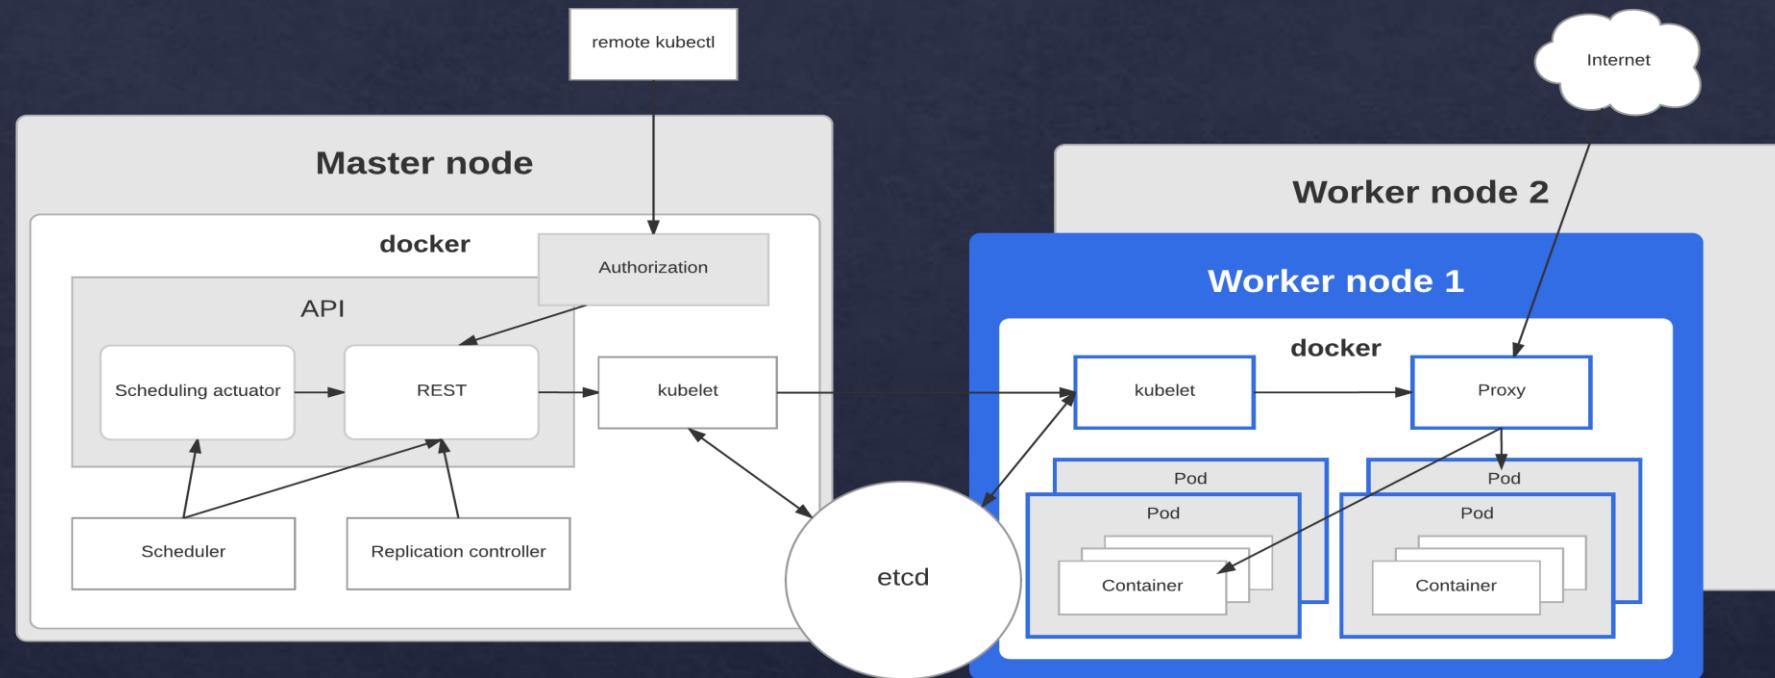
- Kubernetes is a perfect match for microservices
- With Kubernetes, we can create a Deployment to:
 - rollout
 - roll back

Selected containers...

- Kubernetes provides an optional horizontal pod auto-scaling feature
 - by resource
 - custom metrics
 - operates on the ratio between desired metric value and current metric value

Kubernetes Components

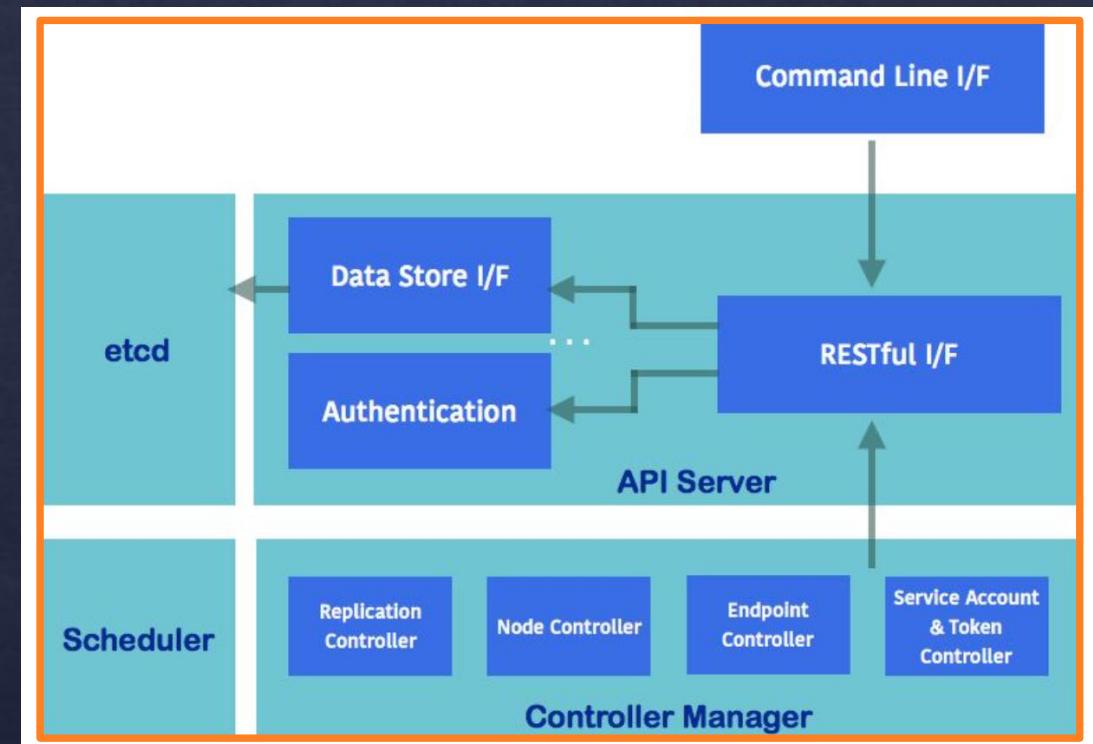
- Kubernetes includes 2 major players



- Masters
 - Is the heart of Kubernetes
 - Controls and schedules all the activities in the cluster
- Nodes
 - Nodes are the workers that run our containers

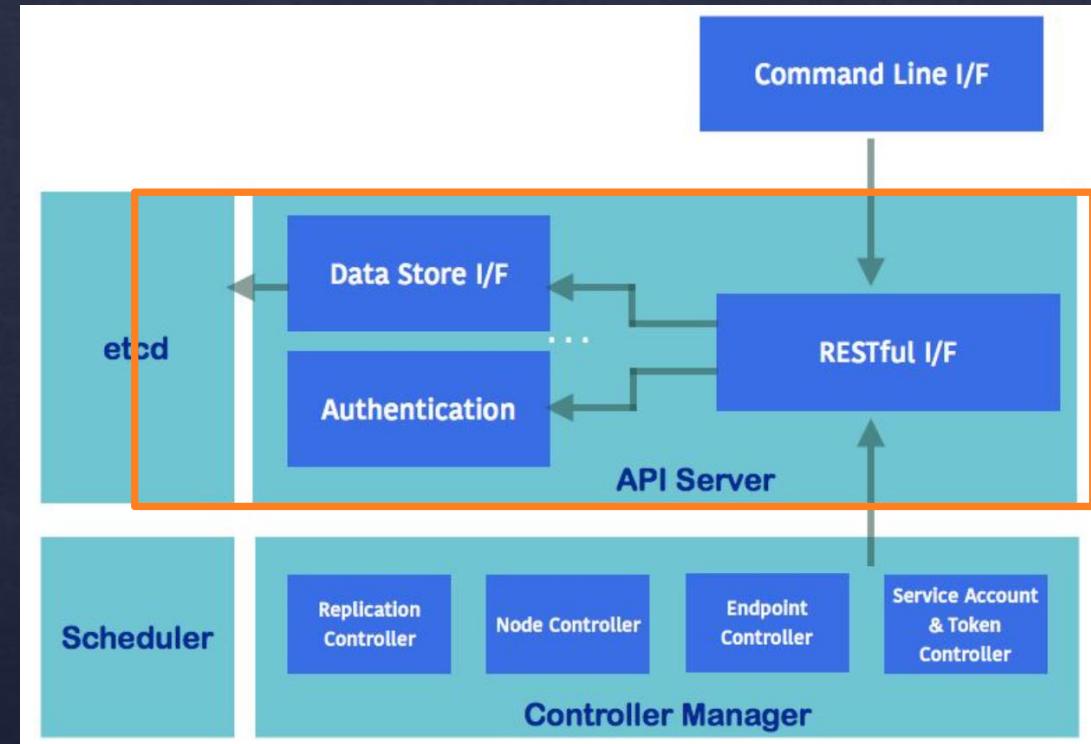
Master Components

- Includes
 1. API server
 2. Controller Manager
 3. Scheduler
 4. Etcd
- All components can run on different hosts with clustering
 - However, from a learning perspective, we'll make all the components run on the same node.

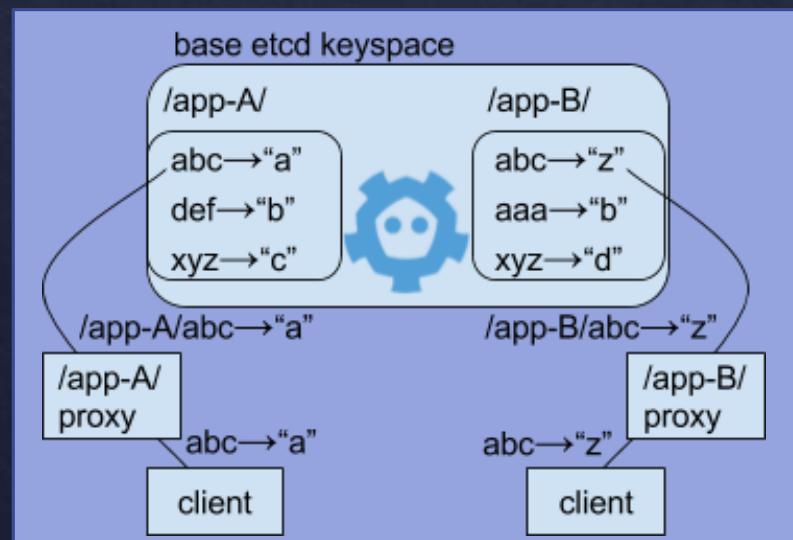


API server (kube-apiserver)

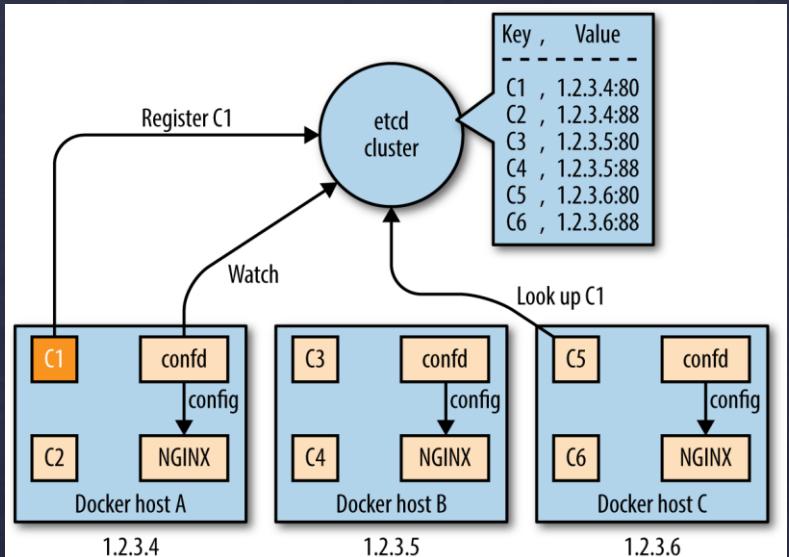
- Provides an HTTP/HTTPS server
 - provides a RESTful API for all the components in the Kubernetes master
- API server **reads and updates etcd**



- etcd is an open source distributed key-value store
- Kubernetes is distributed
 - So, it needs a distributed database



- Kubernetes uses etcd as a key-value database store for
 1. Storing the **configuration** of the Kubernetes cluster in etcd.
 2. Storing the *actual* state of the system and the *desired* state
- Any node crashing or process dying causes values in etcd to be changed.



Controller Manager (kube-controller-manager)

1. Ensures all the ReplicationControllers == Desired container

amount

2. Node Controller Manager responds when the nodes go down,

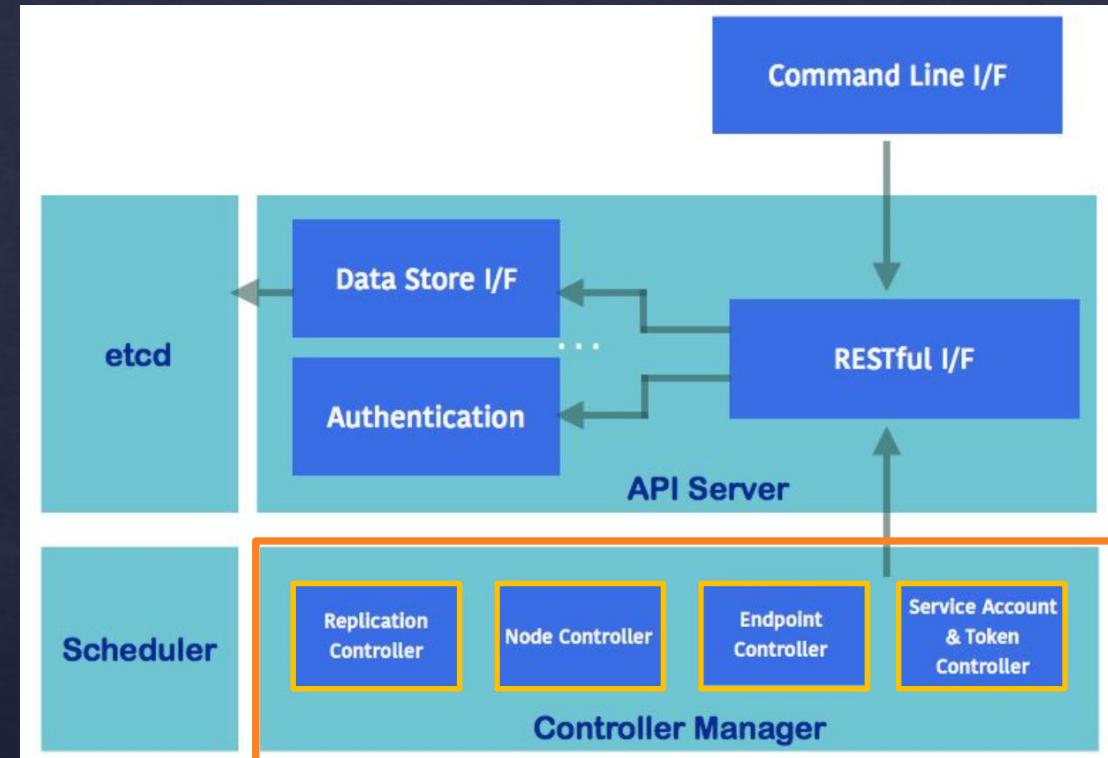
it will then evict the pods

3. Endpoint Controller is used to associate the relationship

between services and pods

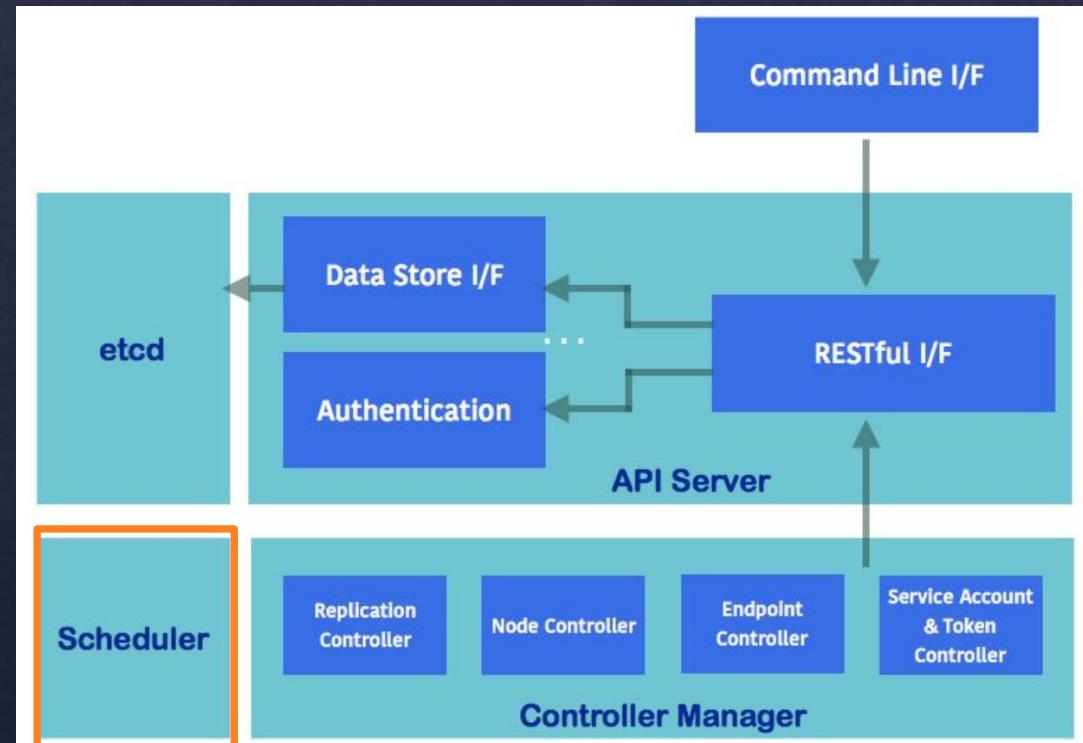
4. Service Account and Token Controller are used to control

default account and API access tokens.



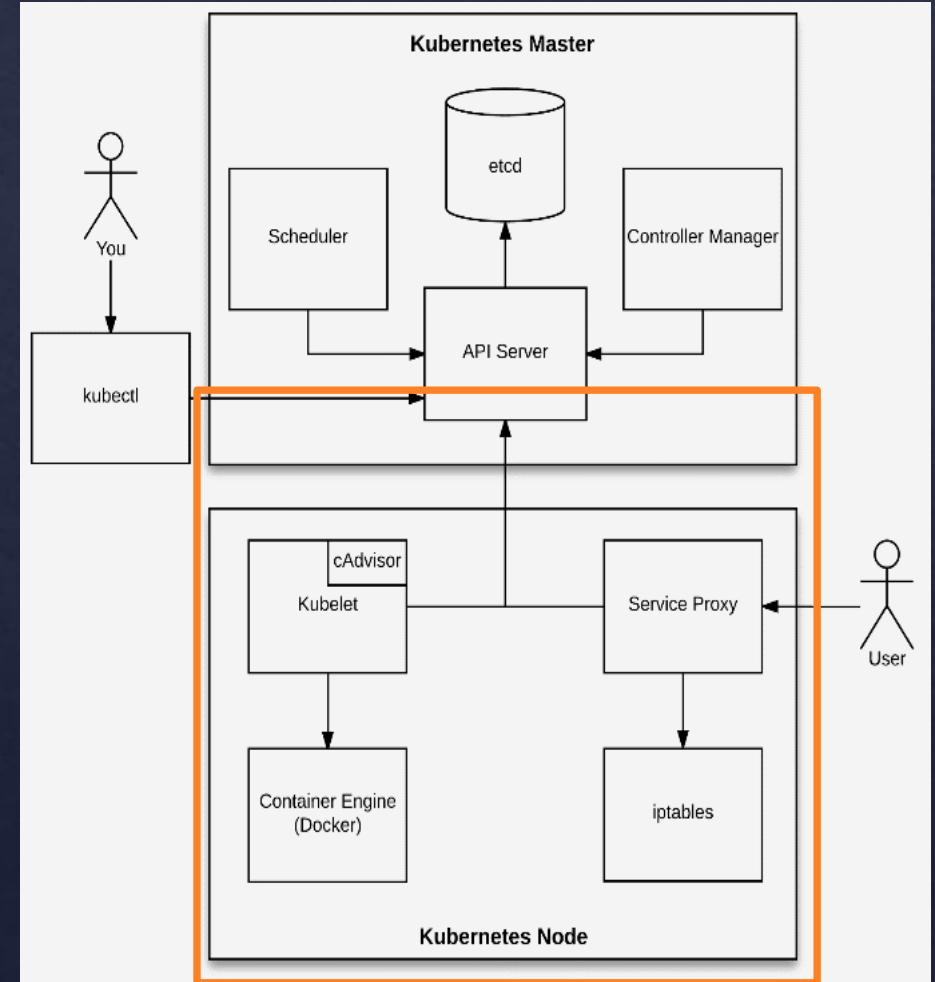
Scheduler (kube-scheduler)

- Decides which node is suitable for pods to run on according to
 - resource capacity
 - balance of the resource utilization on the node
- It also considers spreading the pods in the same set to different nodes.



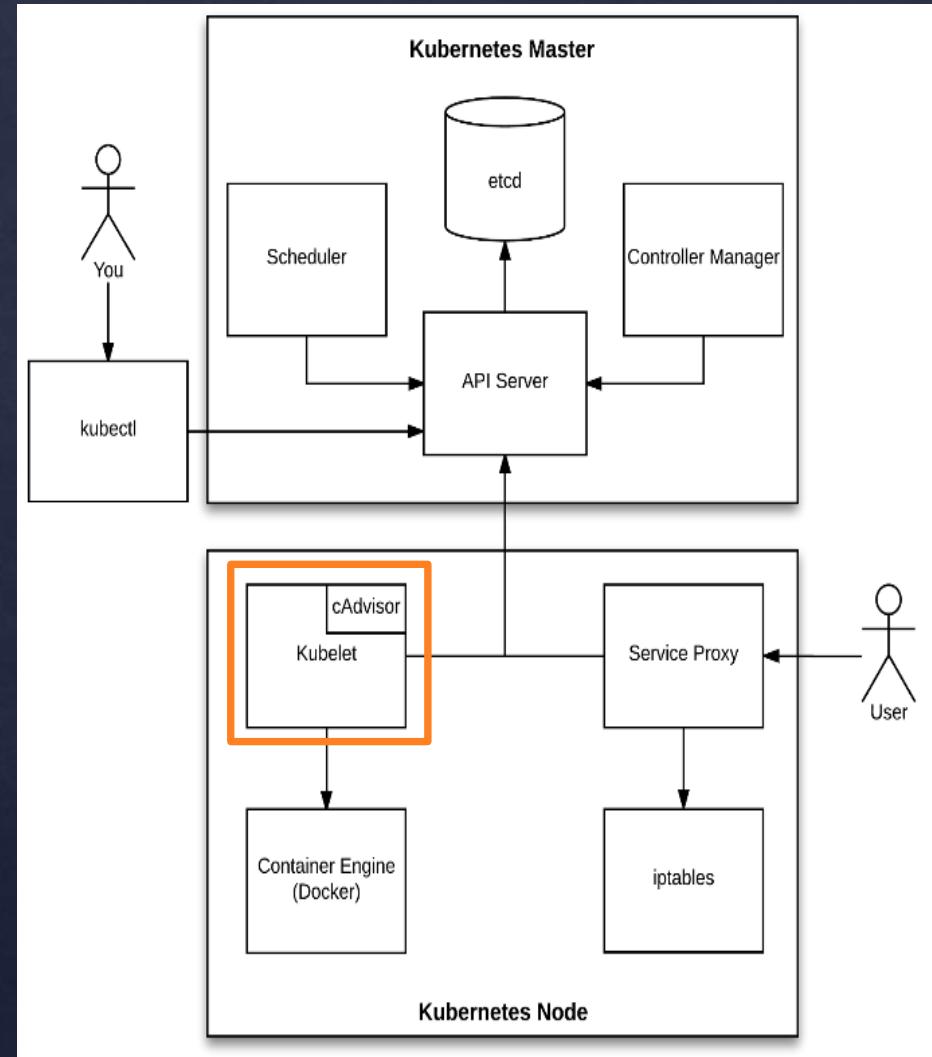
Node Components

- Node components need to be provisioned and run on every node, in order to:
 - report the runtime status of the pod to the master.



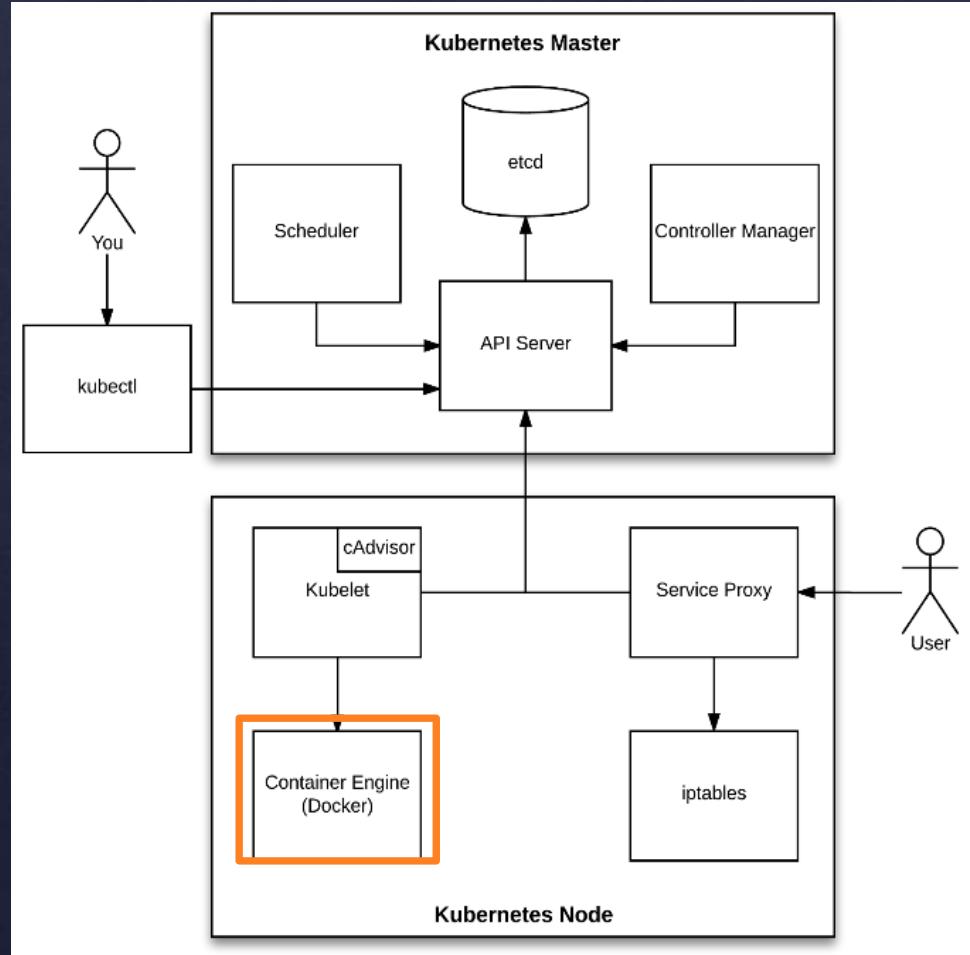
Kubelet

- Is a major process in the nodes
 - Reports node activities to kube-apiserver periodically, such as:
 - pod health
 - node health
- It runs containers via container runtimes, such as Docker or rkt.



Docker Engine

- Kubernetes uses Docker as a default container engine.

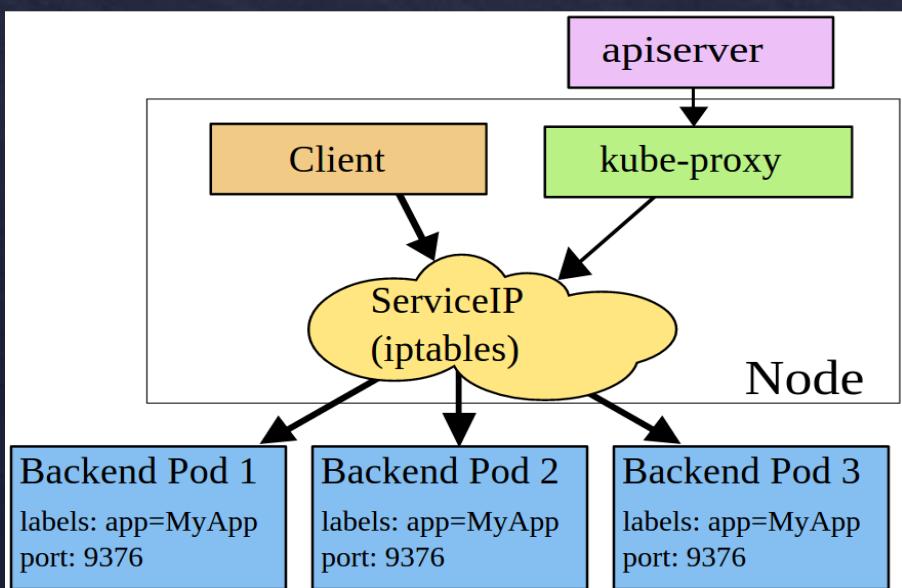
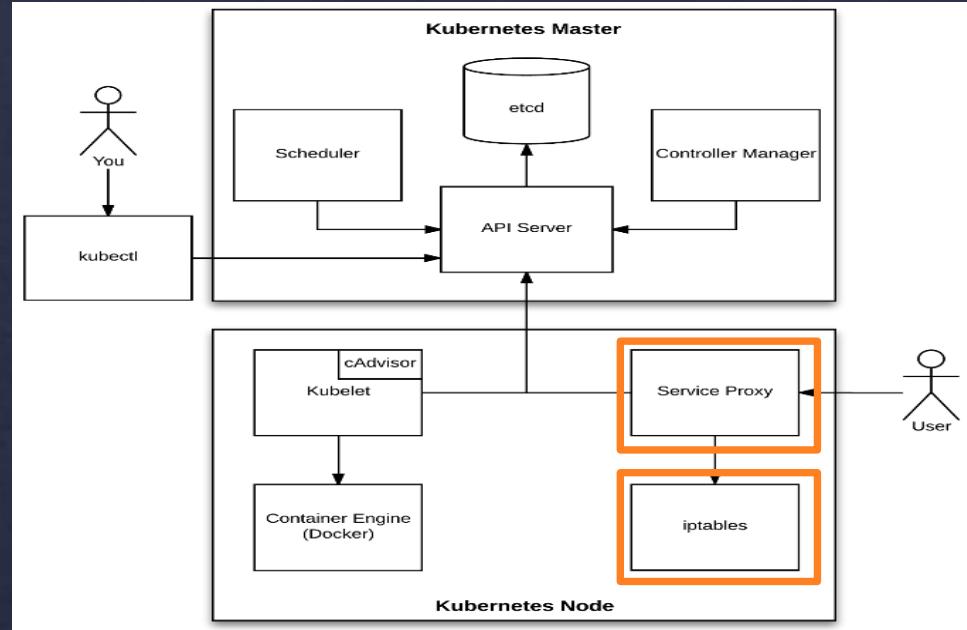


CAdvisor

- The Kubelet ships with built-in support for [cAdvisor](#)
 - It
 - Collects Processes
 - Aggregates Processes
 - Exports metrics
 - CPU
 - Memory
 - File
 - Network Usage

Proxy (kube-proxy)

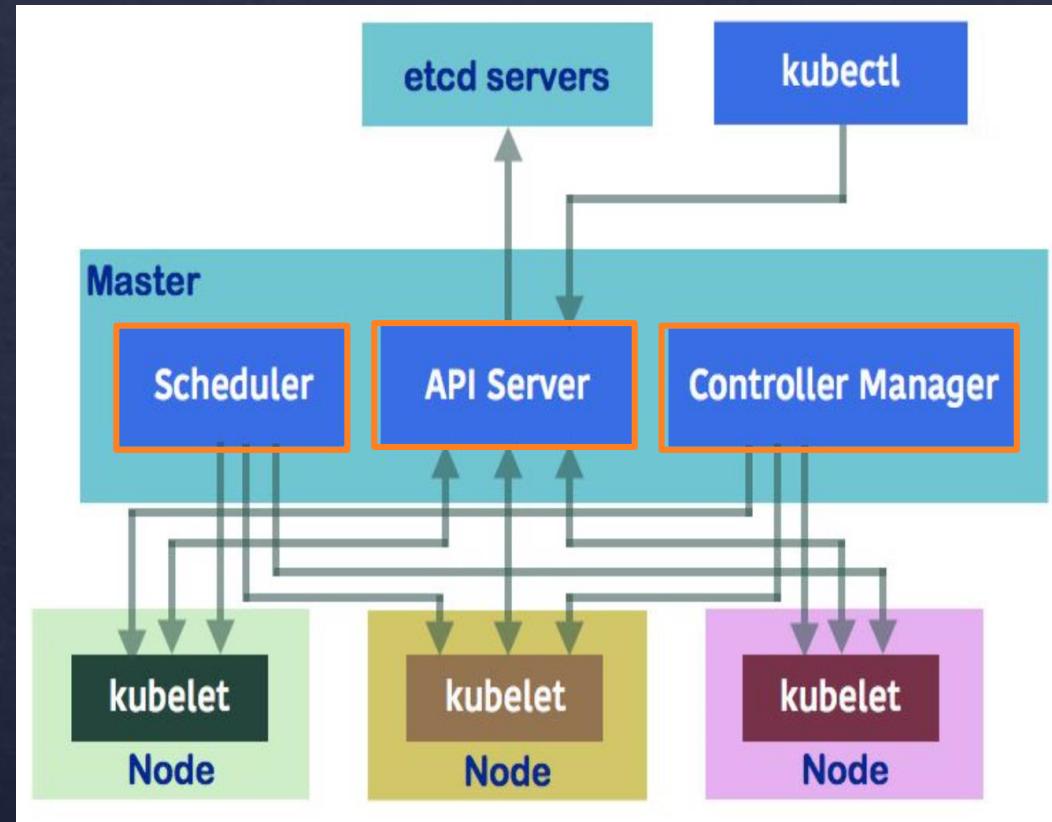
- Handles the routing between pod load-balancer and pods
 - (a.k.a. **service**)
- It also provides the routing from outside to service



- There are **two proxy modes**
 - **Userspace**
 - Large overhead by switching kernel space and user space
 - **Iptables**
 - Default proxy mode
 - It changes iptables **NAT** in Linux to achieve routing TCP and UDP packets across all containers.

Interaction between Kubernetes master and nodes

- **Scheduler** determines which node should be assigned in order to :
 - Run pods
- **Controller Manager** monitors the running tasks and responds if any undesired state occurs
- The **API server** responds to the request after:
 - The client uses **kubectl** to send requests to the API server
 - fetches the logs from pods by kubelet
 - pushes the **object** information from etcd
 - pulls the **object** information from etcd



Activities

Opera



Speed Dial



Enter search or web address



Search the web



+ Add a site

Suggestions



mail.google.com



askubuntu.com

ccm
net

ccm.net

extensions
gnome.org

extensions.gnome.org...

itsfoss
com

itsfoss.com

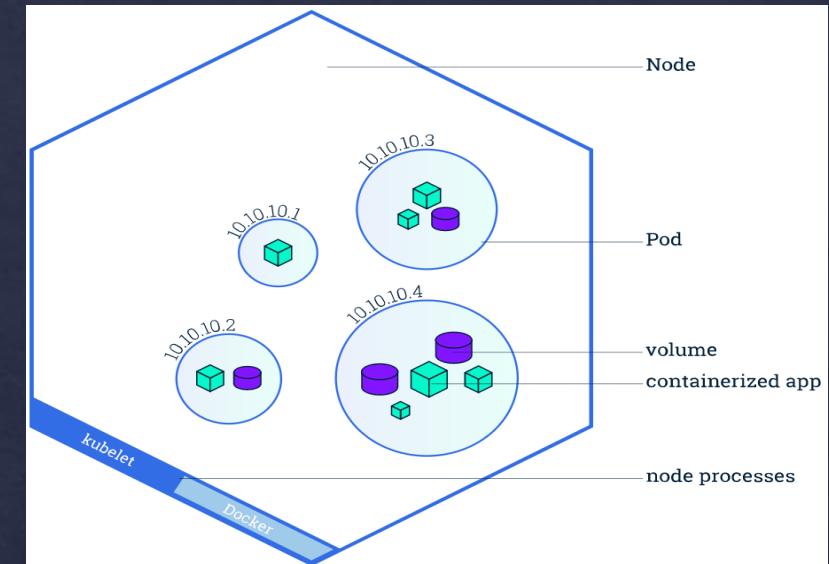
bbs
archlinux
org

bbs.archlinux.org

Objects and Workloads

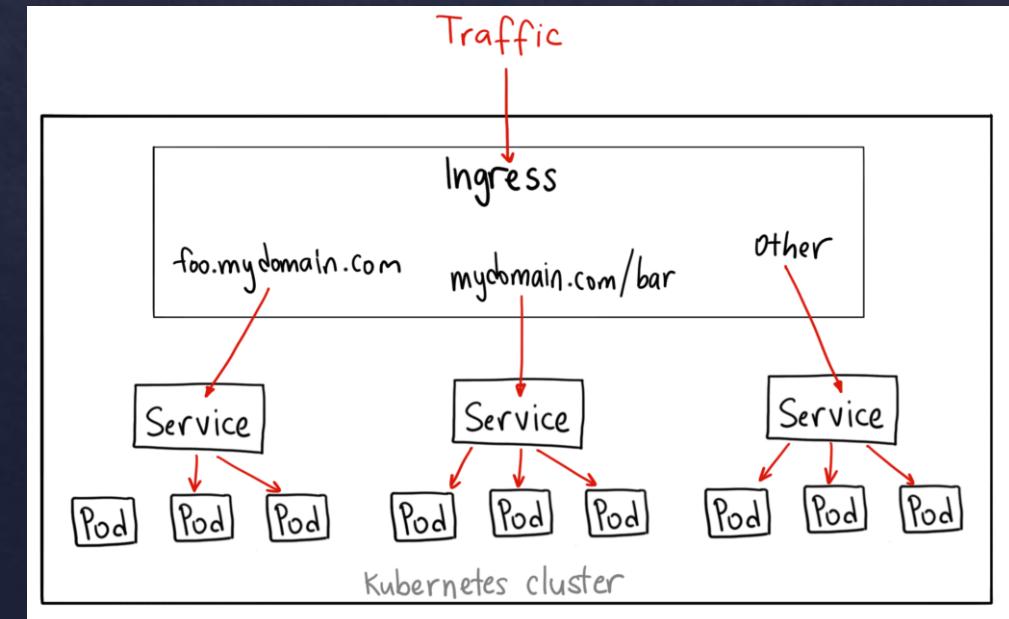
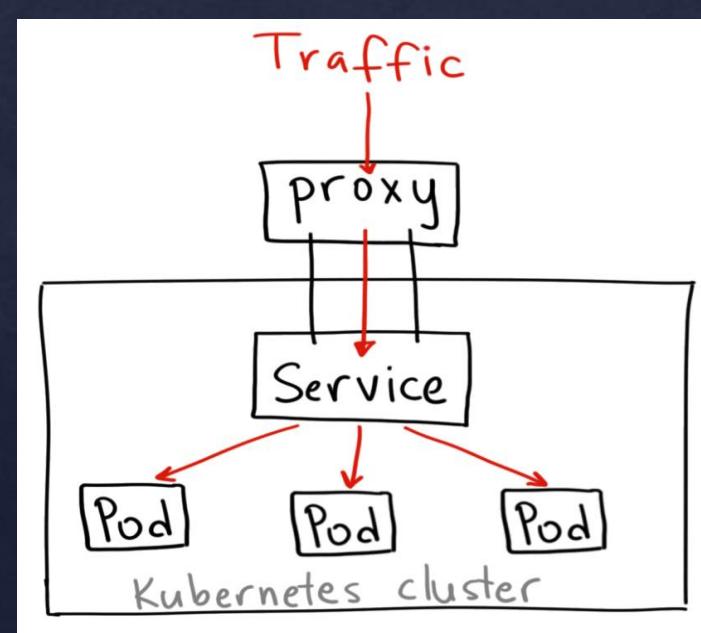
- Nods

- A node is a worker machine in Kubernetes
 - may be a VM/physical machine
 - ✓ depending on the cluster
- Each node contains the services necessary to run pods
- is managed by the master components



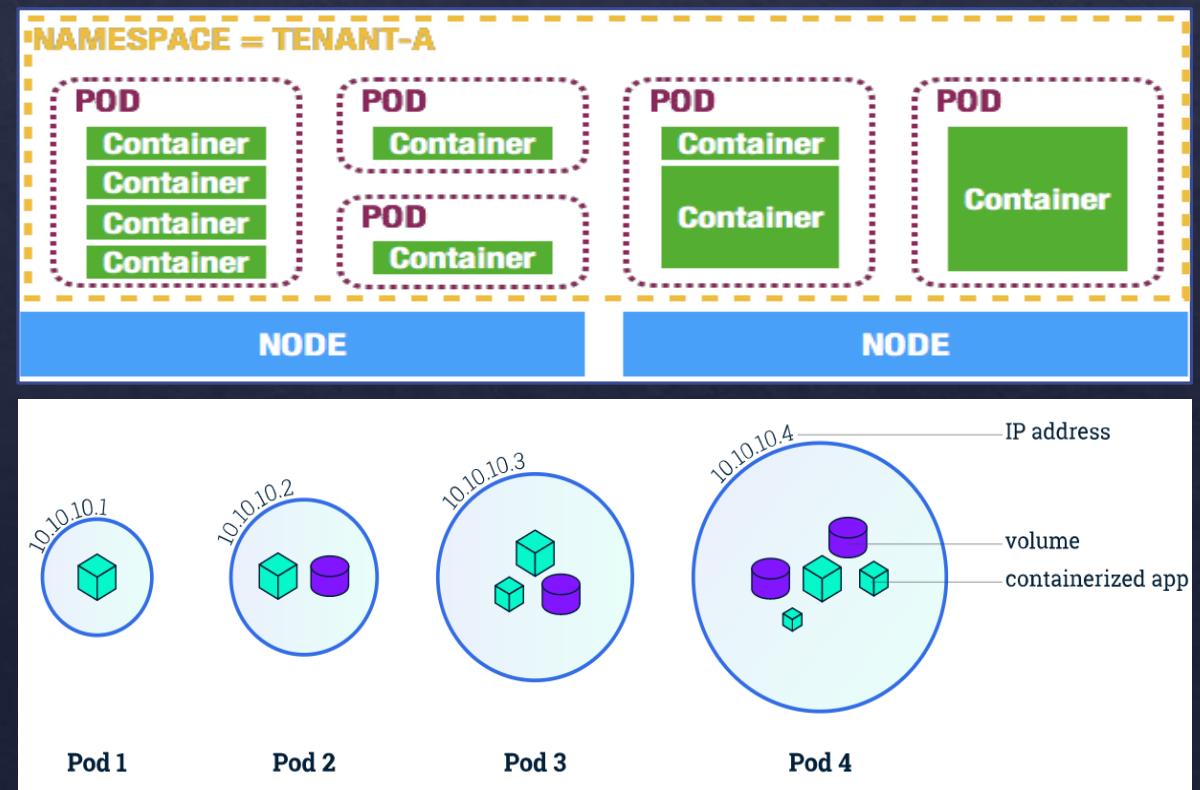
- Services

- Container runtime
- Kubelet
- Kube-Proxy



Objects and Workloads

- Pod
 - Most basic unit
 - One or more **Tightly Coupled Containers** are encapsulated in an object called a Pod.
 - Works as a **single, monolithic application** to best **conceptualize** how the cluster will manage the pod's resources and scheduling.
- Features
 - Share a life cycle
 - Should always be scheduled on the same node
 - managed entirely as a unit
 - share their
 - ✓ Environment
 - ✓ Volumes
 - ✓ IP Space



Objects and Workloads

- Jobs
 - Creates one or more Pods and ensures that a specified number of them successfully terminate
 - A simple case is to create one Job object in order to reliably run one Pod to completion

```
controllers/job.yaml
```

```
apiVersion: batch/v1
kind: Job
metadata:
  name: pi
spec:
  template:
    spec:
      containers:
        - name: pi
          image: perl
          command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
          restartPolicy: Never
        backoffLimit: 4
```

```
kubectl apply -f https://k8s.io/examples/controllers/job.yaml
```

```
job "pi" created
```

```
$ kubectl logs $pods
3.14159265358979323846264338327950288419716939937510582097494459230781640628620899862
```

Jobs

```
apiVersion: batch/v1
kind: Job
metadata:
  name: pi
spec:
  template:
    spec:
      containers:
        - name: pi
          image: perl
          command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
      restartPolicy: Never
  backoffLimit: 4
```

```
kubectl apply -f https://k8s.io/examples/controllers/job.yaml
```

```
job "pi" created
```

controllers/job.yaml

kubectl describe jobs/pi

```
Name:           pi
Namespace:      default
Selector:       controller-uid=b1db589a-2c8d-11e6-b324-0209dc45a495
Labels:         controller-uid=b1db589a-2c8d-11e6-b324-0209dc45a495
                job-name=pi
```

```
Annotations:    <none>
Parallelism:   1
Completions:   1
Start Time:    Tue, 07 Jun 2016 10:56:16 +0200
Pods Statuses: 0 Running / 1 Succeeded / 0 Failed
Pod Template:
  Labels:       controller-uid=b1db589a-2c8d-11e6-b324-0209dc45a495
                job-name=pi
```

```
Containers:
  pi:
    Image:      perl
    Port:       0
    Command:
      perl
      -Mbignum=bpi
      -wle
      print bpi(2000)
```

```
Environment:   <none>
Mounts:        <none>
Volumes:       <none>
```

FirstSeen	LastSeen	Count	From	SubobjectPath	Type	Reason
1m	1m	1	{job-controller }		Normal	SuccessfulCreate

```
pods=$(kubectl get pods --selector=job-name=pi --output=jsonpath='{.items[*].metadata.name}')
echo $pods
```

```
$ kubectl logs $pods
3.14159265358979323846264338327950288419716939937510582097494459230781640628620899862
```

Activities Terminal



kubernetes-u18@ubuntu:~



File Edit View Search Terminal Help

kubernetes-u18@ubuntu:~\$

I



❖ Jobs

Objects and Workloads

- Replication Controllers
 - Is an object that defines a pod template
 - An easy way to **distribute load**
 - An easy way to **increase availability**
 - Control parameters to scale identical replicas of a pod horizontally
 - **EX.** increase the desired instance count from 3 to 4
 - Ensures that the number of pods deployed in the cluster **==** The number of pods in its configuration
 - If a pod or underlying host fails, the controller will start new pods to compensate
 - ✓ Each replication controller has a **desired state**
 - Knows how to create new pods as needed
 - A template that resembles a pod definition is embedded within the replication controller configuration
 - Replication controllers can also perform rolling updates to roll over a set of pods to a new version

Replication Controllers

- Makes use of label queries

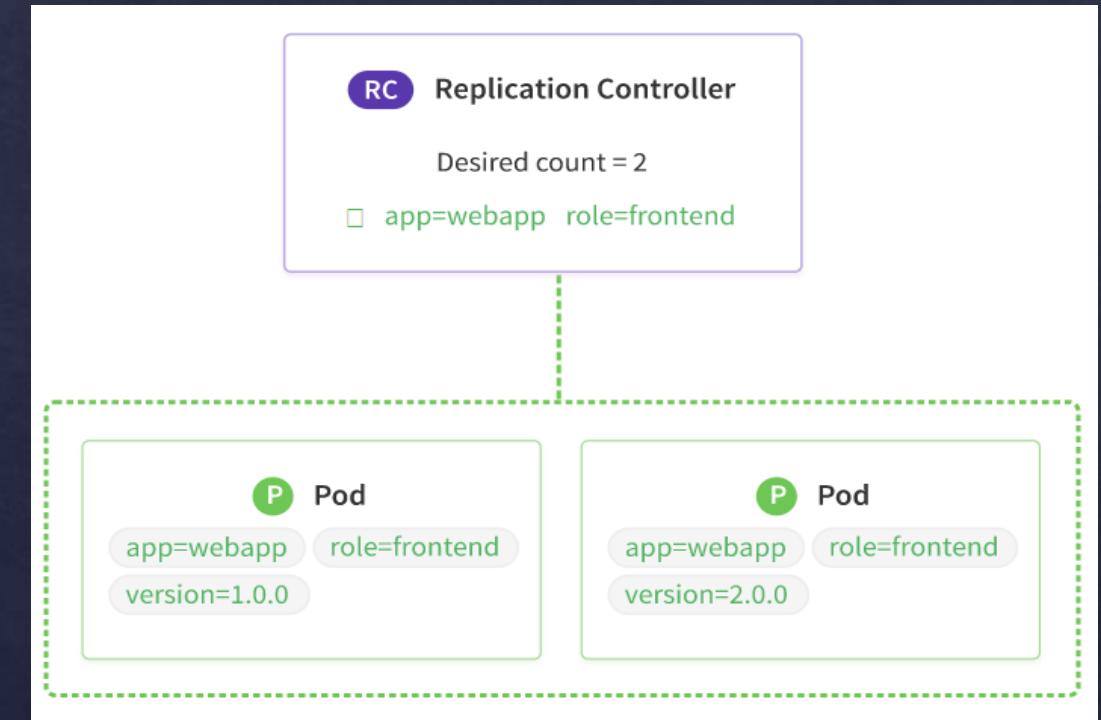
```
controllers/replication.yaml □  
  
apiVersion: v1  
kind: ReplicationController  
metadata:  
  name: nginx  
spec:  
  replicas: 3  
  selector:  
    app: nginx  
  template:  
    metadata:  
      name: nginx  
      labels:  
        app: nginx  
    spec:  
      containers:  
      - name: nginx  
        image: nginx  
        ports:  
        - containerPort: 80
```



```
kubectl apply -f https://k8s.io/examples/controllers/replication.yaml  
replicationcontroller/nginx created
```



```
kubectl describe replicationcontrollers/nginx
```



Replication Sets

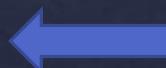
- Ensures that a specified number of pod replicas are running at any given time
- The link a ReplicaSet has to its Pods is via the Pods' `metadata.ownerReferences` field
 - Specifies what resources the current object is owned by
 - All Pods acquired by a ReplicaSet have their owning ReplicaSet's identifying information within their `ownerReferences` field

```
kubectl apply -f https://kubernetes.io/examples/controllers/frontend.yaml
```



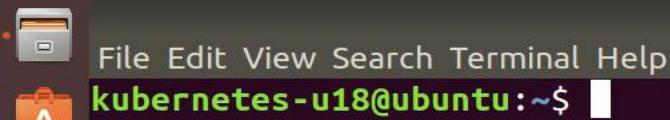
```
kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
frontend	3	3	3	6s



```
controllers/frontend.yaml □
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: frontend
  labels:
    app: guestbook
    tier: frontend
spec:
  # modify replicas according to your case
  replicas: 3
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v3
```

Activities Terminal



kubernetes-u18@ubuntu:~

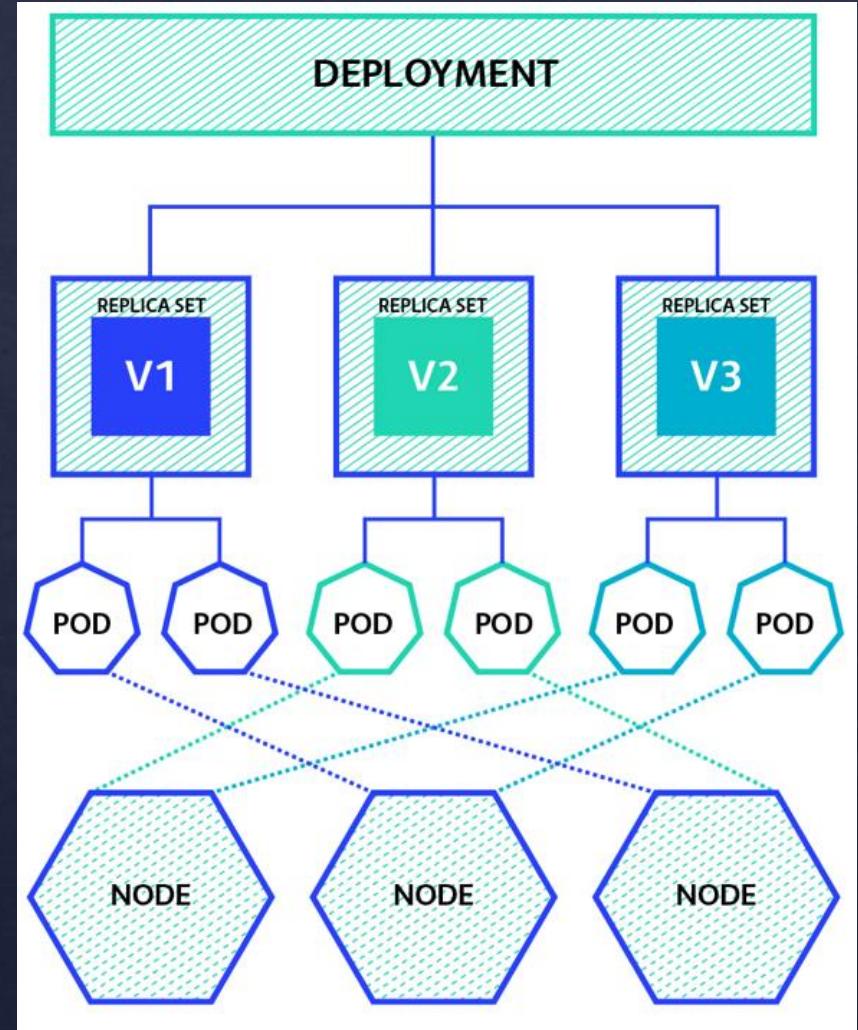
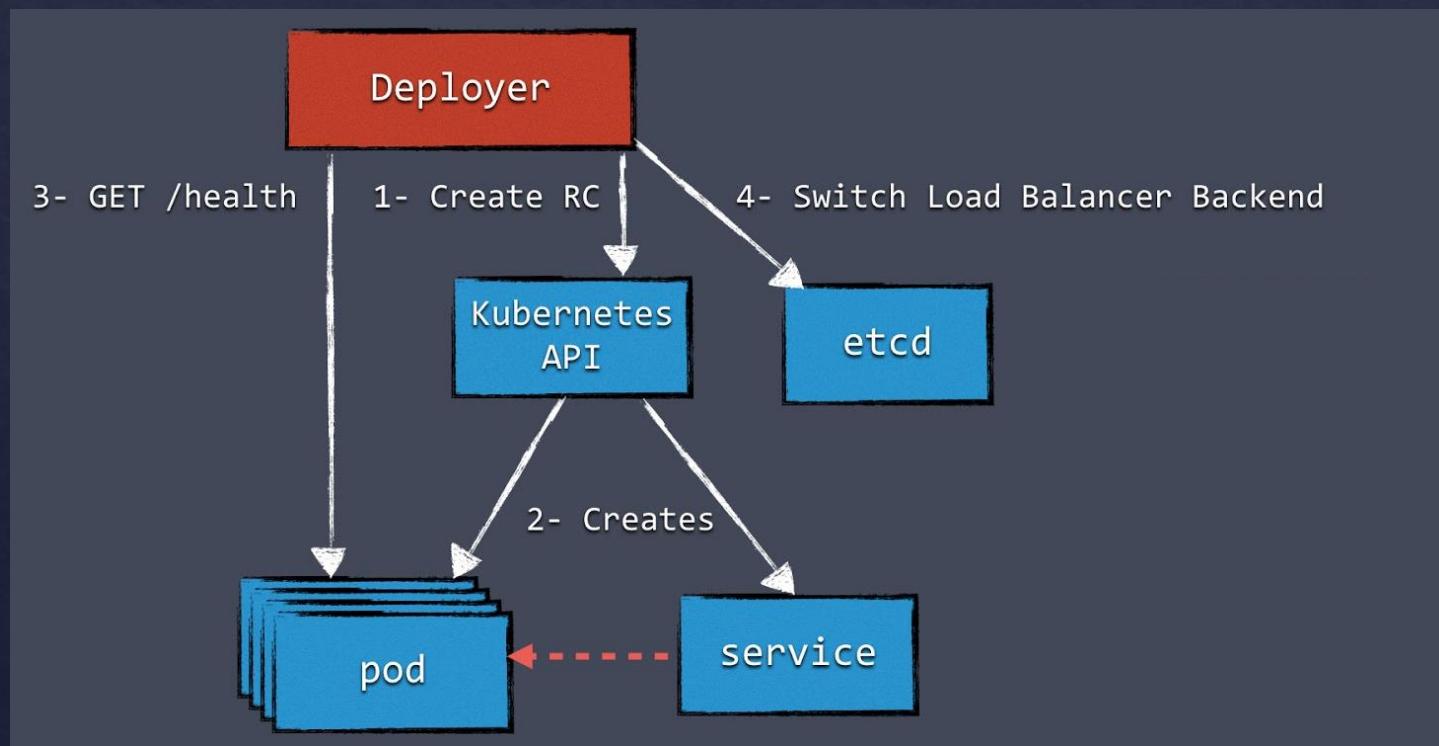
File Edit View Search Terminal Help
kubernetes-u18@ubuntu:~\$

Objects and Workloads

- Replication Controllers and Replication Sets
 - Managing groups of identical, replicated pods instead of single pods
 - These are created from **pod templates**
 - Are the same **except** the “**Roll Over**” Feature
 - **Lack** some fine grained life cycle management capabilities found in more complex objects.

Objects and Workloads

- Deployments
 - One of the most common workloads to directly create and manage.
 - Use replication sets as a building block.



Deployments

- One of the most common workloads to directly create and manage.
- Use replication sets as a building block.

```
kubectl apply -f https://k8s.io/examples/controllers/nginx-deployment.yaml
```

```
kubectl get deployments
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
nginx-deployment	3	0	0	0	1s

```
kubectl rollout status deployment.v1.apps/nginx-deployment .
```

```
Waiting for rollout to finish: 2 out of 3 new replicas have been updated...
deployment.apps/nginx-deployment successfully rolled out
```

```
kubectl get deployments
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
nginx-deployment	3	3	3	3	18s

```
kubectl --record deployment.apps/nginx-deployment set image deployment.v1.apps/nginx-deployment nginx=nginx:1.9.1
```

```
image updated
```

```
controllers/nginx-deployment.yaml
```

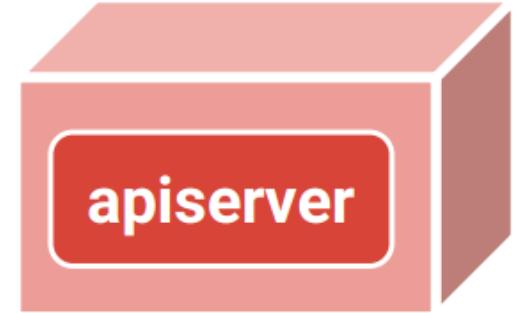
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```

Deployment

Problem: Manual rollout of applications (and changes)

Today:

- kubectl rolling-update
- sed (?!?)

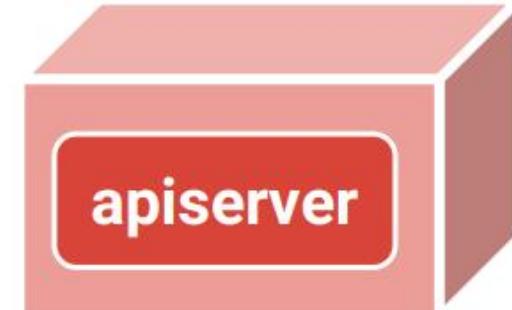


Deployment API

Solution: Deployment API

```
$ kubectl create -f Deployment-1.yml
```

```
metadata:
  name: sweet-webapp
spec:
  replicas: 42
  template:
    spec:
      containers:
        - name: webapp
          image: webapp:2.1.7
          ports:
            - containerPort: 80
```

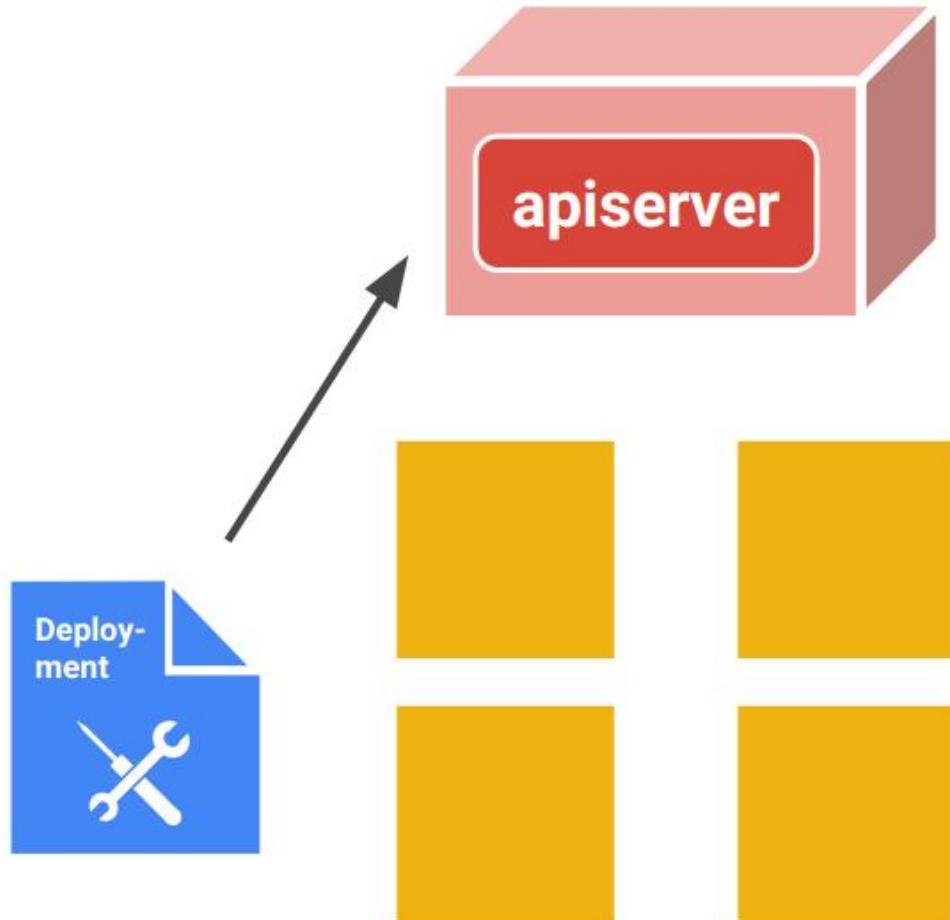


Deployment API

Solution: Deployment API

```
$ kubectl create -f Deployment-1.yml
```

```
metadata:  
  name: sweet-webapp  
spec:  
  replicas: 42  
  template:  
    spec:  
      containers:  
        - name: webapp  
          image: webapp:2.1.7  
          ports:  
            - containerPort: 80
```

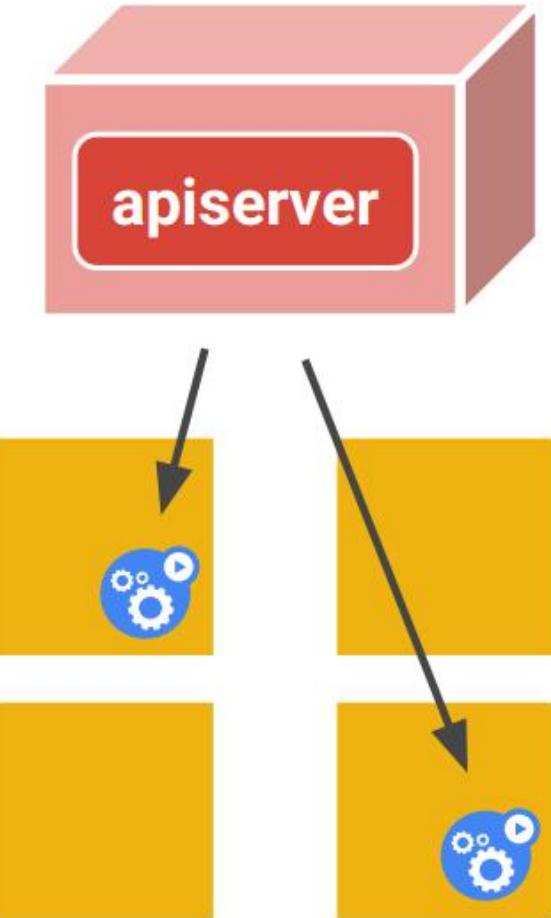


Deployment API

Solution: Deployment API

```
$ kubectl create -f Deployment-1.yml
```

```
metadata:  
  name: sweet-webapp  
spec:  
  replicas: 42  
  template:  
    spec:  
      containers:  
      - name: webapp  
        image: webapp:2.1.7  
        ports:  
        - containerPort: 80
```

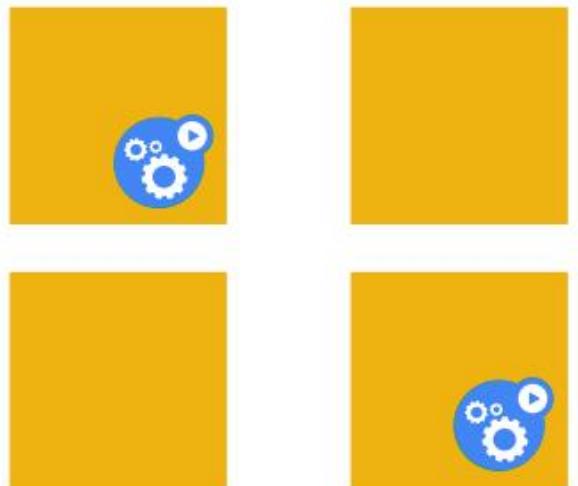
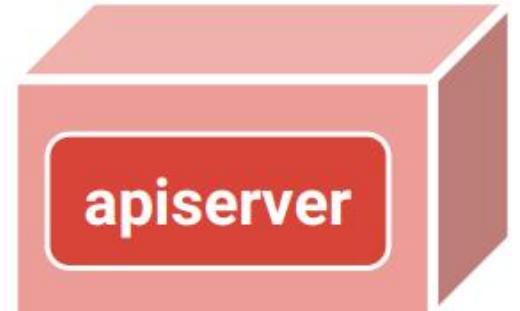


Deployment API

Solution: Deployment API!

```
$ kubectl create -f Deployment-2.yml
```

```
metadata:  
  name: sweet-webapp  
spec:  
  replicas: 42  
  template:  
    spec:  
      containers:  
        - name: webapp  
          image: webapp:2.1.8  
          ports:  
            - containerPort: 80
```

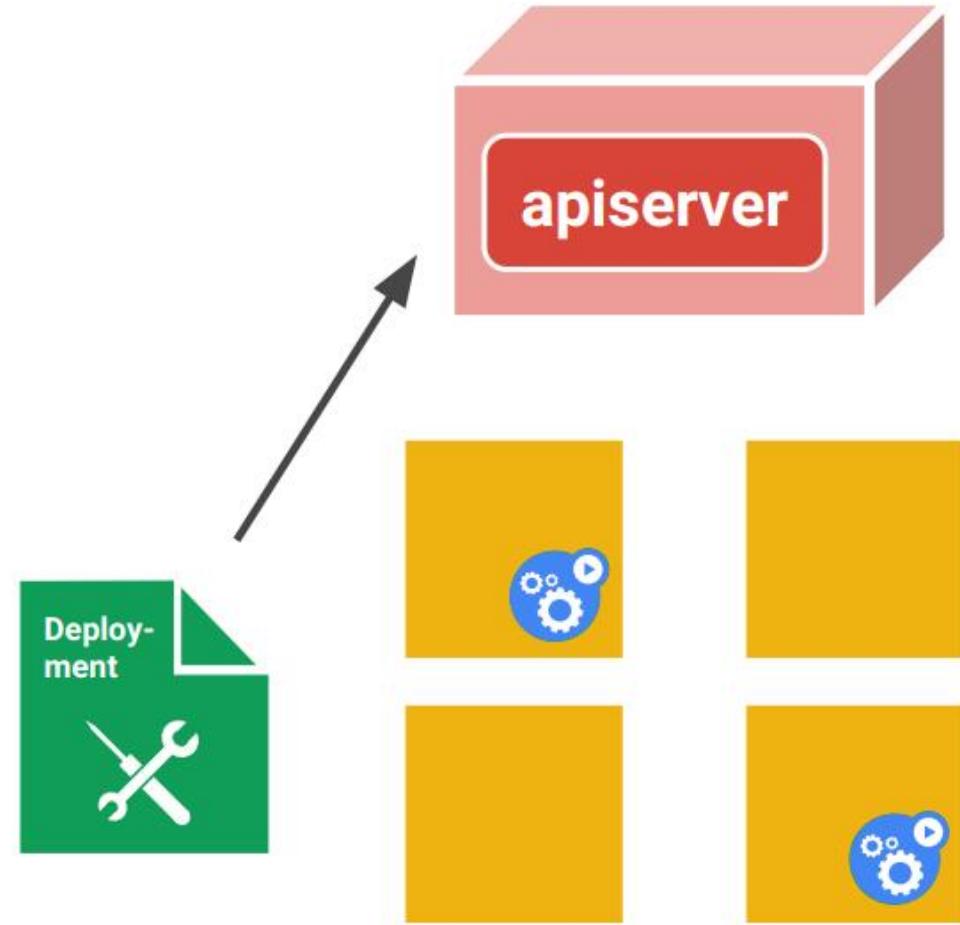


Deployment API

Solution: Deployment API!

```
$ kubectl create -f Deployment-2.yml
```

```
metadata:  
  name: sweet-webapp  
spec:  
  replicas: 42  
  template:  
    spec:  
      containers:  
        - name: webapp  
          image: webapp:2.1.8  
          ports:  
            - containerPort: 80
```

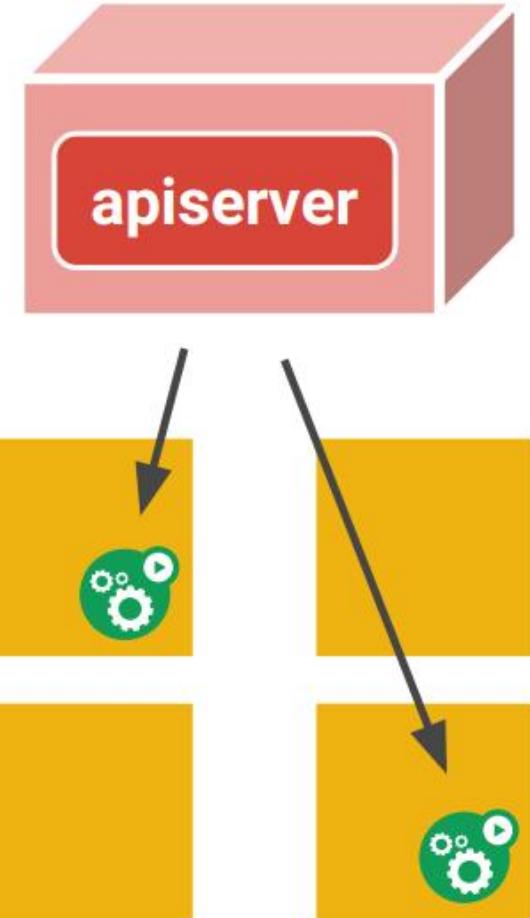


Deployment API

Solution: Deployment API!

```
$ kubectl create -f Deployment-2.yml
```

```
metadata:  
  name: sweet-webapp  
spec:  
  replicas: 42  
  template:  
    spec:  
      containers:  
        - name: webapp  
          image: webapp:2.1.8  
          ports:  
            - containerPort: 80
```

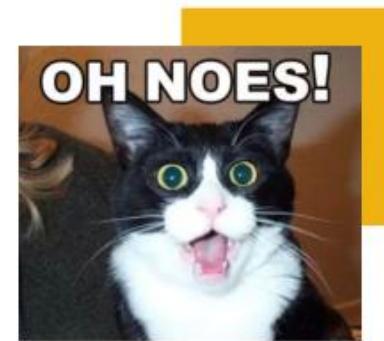


Deployment API

Solution: Deployment API!

```
$ kubectl create -f Deployment-2.yml
```

```
metadata:  
  name: sweet-webapp  
spec:  
  replicas: 42  
  template:  
    spec:  
      containers:  
        - name: webapp  
          image: webapp:2.1.8  
          ports:  
            - containerPort: 80
```

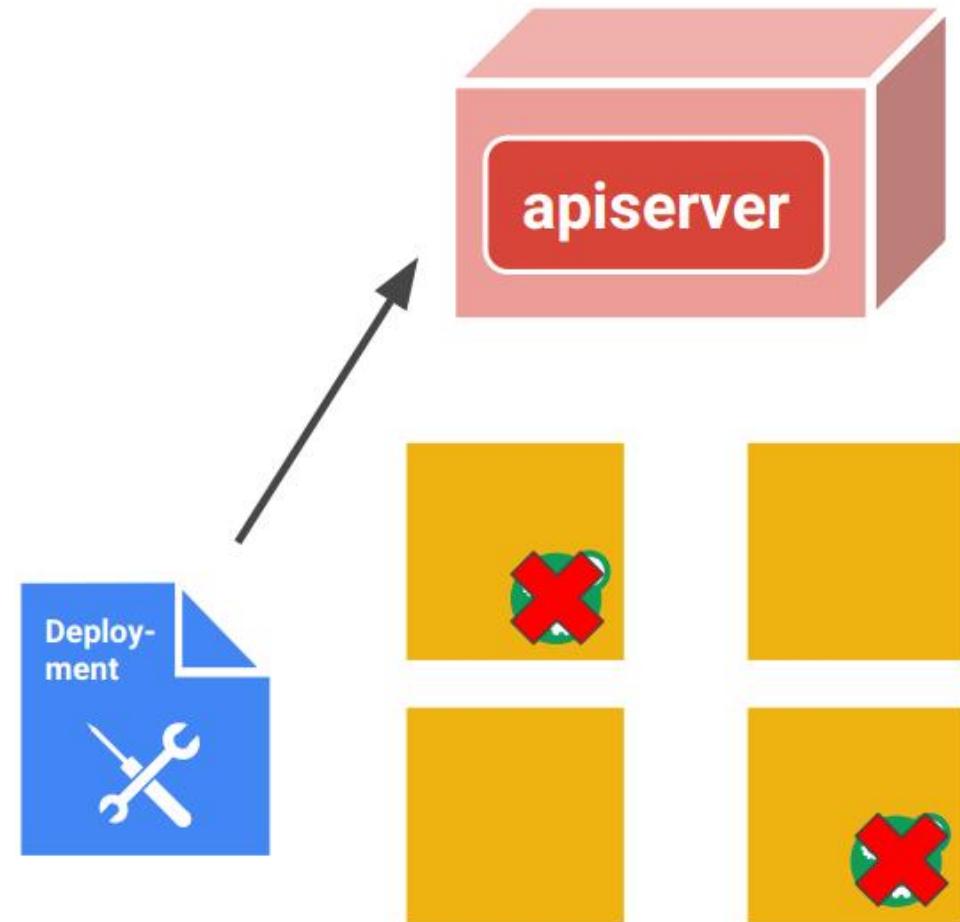


Deployment API

Solution: Deployment API!

```
$ kubectl create -f Deployment-1.yml
```

```
metadata:  
  name: sweet-webapp  
spec:  
  replicas: 42  
  template:  
    spec:  
      containers:  
        - name: webapp  
          image: webapp:2.1.7  
          ports:  
            - containerPort: 80
```



Deployment API

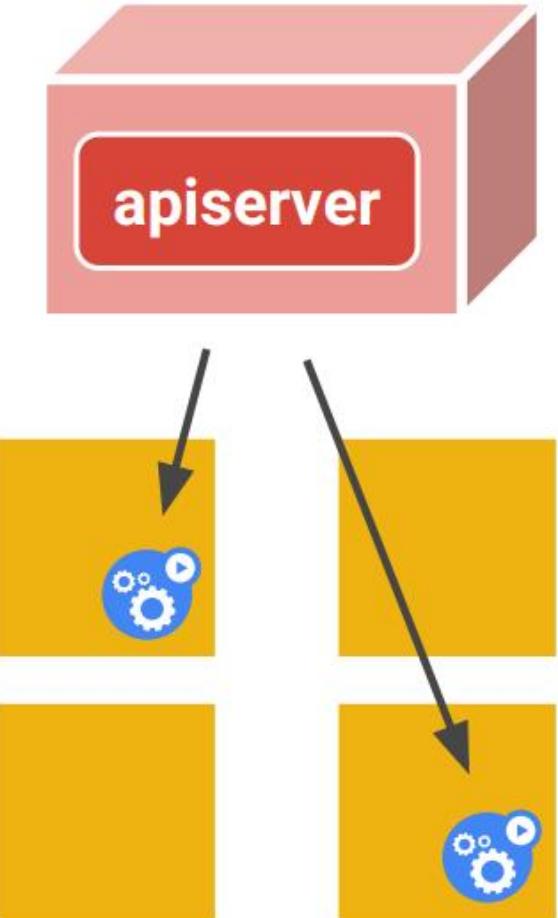
Solution: Deployment API!

```
$ kubectl create -f Deployment-1.yml
```

```
metadata:  
  name: sweet-webapp  
spec:  
  replicas: 42  
  template:  
    spec:  
      containers:  
        - name: webapp  
          image: webapp:2.1.7  
          ports:  
            - containerPort: 80
```



Excellent.



Activities Terminal



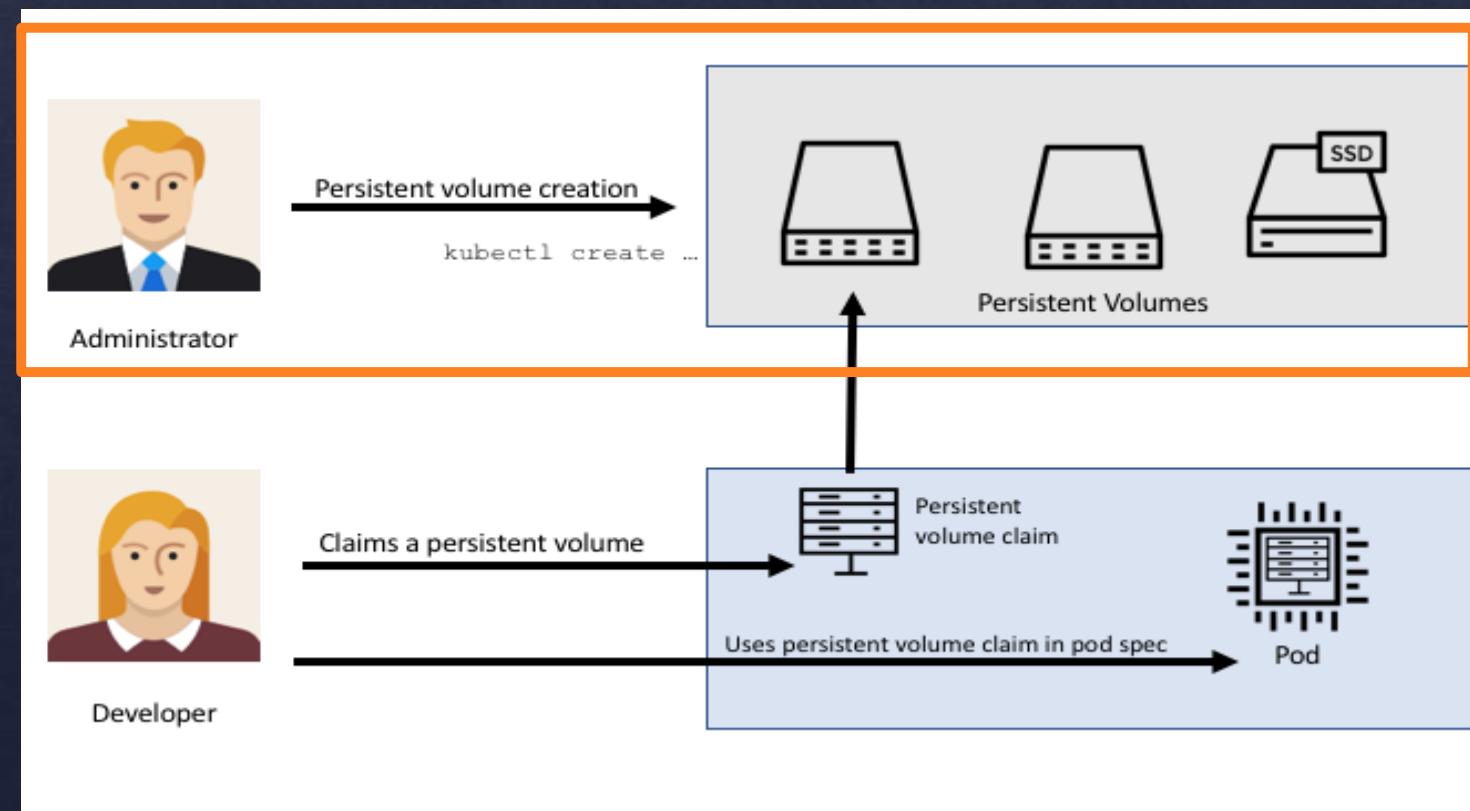
kubernetes-u18@ubuntu:~

File Edit View Search Terminal Help
kubernetes-u18@ubuntu:~\$

Objects and Workloads

- Persistent Volume(PV)
 - Piece of storage in the cluster that has been provisioned by an administrator
 - This API object captures the details of the implementation of the storage, be that NFS, iSCSI, or a cloud-provider-specific storage system.

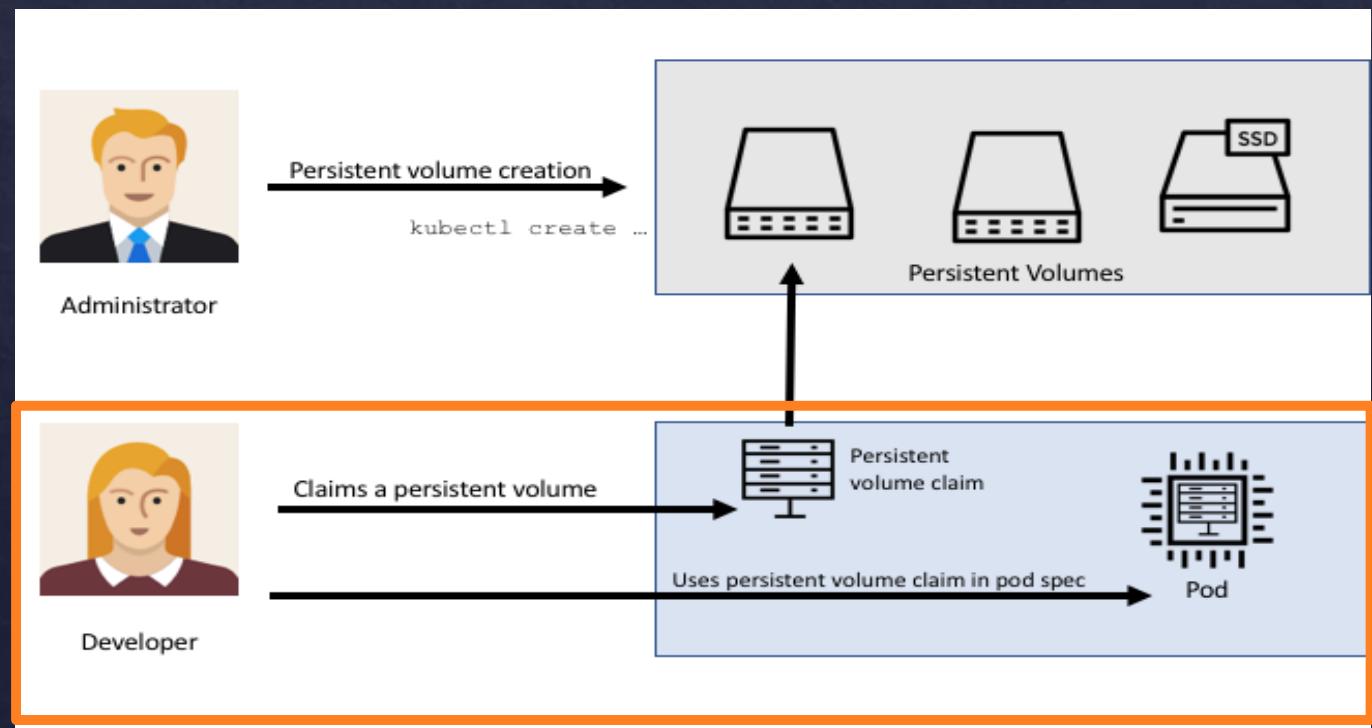
```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv0003
spec:
  capacity:
    storage: 5Gi
  volumeMode: Filesystem
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Recycle
  storageClassName: slow
  mountOptions:
    - hard
    - nfsvers=4.1
  nfs:
    path: /tmp
    server: 172.17.0.2
```



Objects and Workloads

- Persistent Volume Claim(PVC)
 - It is similar to a pod Except
 1. Pods consume **node resources** While PVCs consume **PV resources**
 2. Pods can request specific levels of resources (CPU and Memory)(**Namespace**).
 3. Claims can request specific size and access modes (can be mounted once read/write or many times read-only).

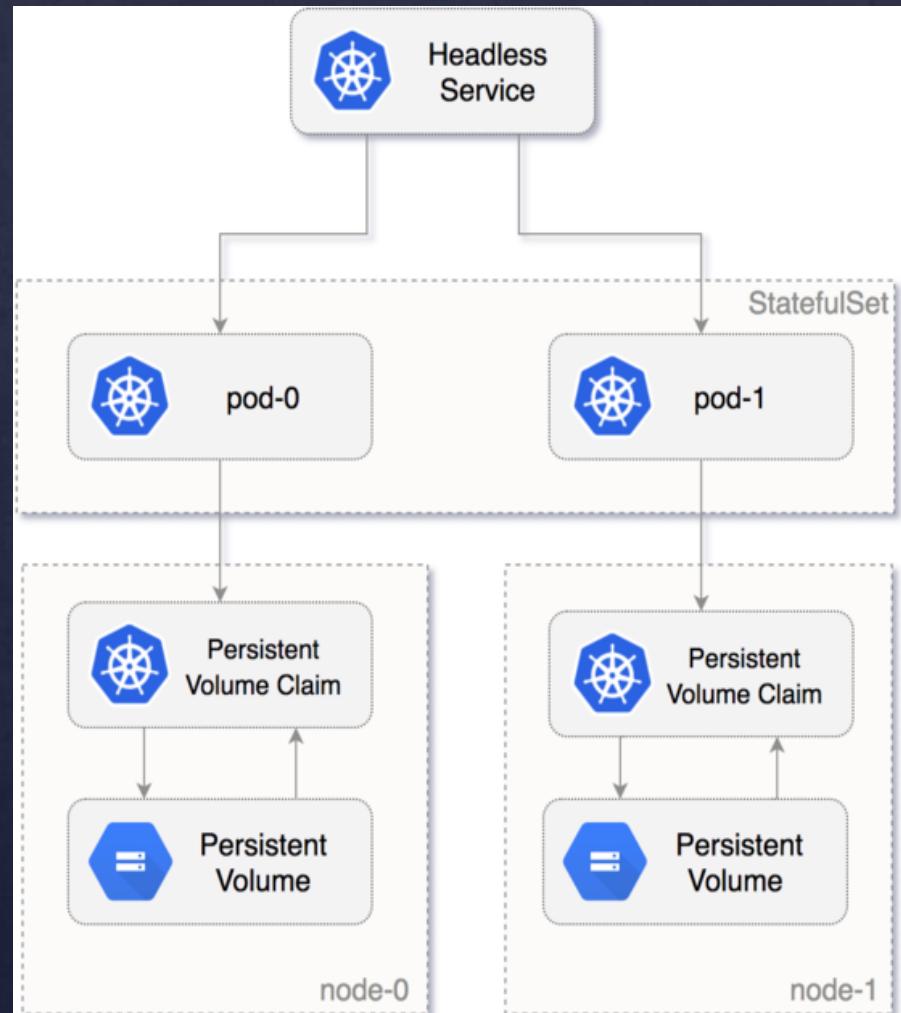
```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: myclaim
spec:
  accessModes:
    - ReadWriteOnce
  volumeMode: Filesystem
  resources:
    requests:
      storage: 8Gi
  storageClassName: slow
  selector:
    matchLabels:
      release: "stable"
    matchExpressions:
      - {key: environment, operator: In, values: [dev]}
```



Objects and Workloads

- StatefulSets

- Are specialized **pod controllers**
- Often associated with **data-oriented applications**
 - Like databases,
 - ✓ which need access to the same volumes even if rescheduled to a new node.
- Gives fine-grained control when you have special requirements related to :
 - Stable, unique network identifiers.
 - Stable, persistent storage.
 - Ordered, graceful deployment and scaling.
 - Ordered, automated rolling updates.



StatefulSets

- Update Strategies

- `.spec.updateStrategy.type` is set to `RollingUpdate`
- `.spec.updateStrategy.type` is set to `onDelete`

```
apiVersion: v1
kind: Service
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  ports:
  - port: 80
    name: web
  clusterIP: None
  selector:
    app: nginx
---
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: web
spec:
  selector:
    matchLabels:
      app: nginx # has to match .spec.template.metadata.labels
  serviceName: "nginx"
  replicas: 3 # by default is 1
  template:
    metadata:
      labels:
        app: nginx # has to match .spec.selector.matchLabels
  spec:
    terminationGracePeriodSeconds: 10
    containers:
    - name: nginx
      image: k8s.gcr.io/nginx-slim:0.8
      ports:
      - containerPort: 80
        name: web
      volumeMounts:
      - name: www
        mountPath: /usr/share/nginx/html
  volumeClaimTemplates:
  - metadata:
      name: www
  spec:
    accessModes: [ "ReadWriteOnce" ]
    storageClassName: "my-storage-class"
    resources:
      requests:
        storage: 1Gi
```

Objects and Workloads

- Daemon Sets
 - Another specialized form of pod controller
 - Run a copy of a pod on each node in the cluster (or a subset, if specified)
 - Useful when deploying pods that help {perform maintenance/provide services} for the nodes
 - Example:
 - running a cluster storage daemon, such as glusterd, ceph, on each node
 - running a logs collection daemon on every node, such as fluentd or logstash
 - running a node monitoring daemon on every node, such as
 - ✓ Prometheus Node Exporter
 - ✓ Collectd
 - ✓ SignalFx Agent

Daemon Sets

```
kubectl apply -f https://k8s.io/examples/controllers/daemonset.yaml
```

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: fluentd-elasticsearch
  namespace: kube-system
  labels:
    k8s-app: fluentd-logging
spec:
  selector:
    matchLabels:
      name: fluentd-elasticsearch
  template:
    metadata:
      labels:
        name: fluentd-elasticsearch
    spec:
      tolerations:
        - key: node-role.kubernetes.io/master
          effect: NoSchedule
      containers:
        - name: fluentd-elasticsearch
          image: gcr.io/fluentd-elasticsearch/fluentd:v2.5.1
          resources:
            limits:
              memory: 200Mi
            requests:
              cpu: 100m
              memory: 200Mi
          volumeMounts:
            - name: varlog
              mountPath: /var/log
            - name: varlibdockercontainers
              mountPath: /var/lib/docker/containers
              readOnly: true
      terminationGracePeriodSeconds: 30
  volumes:
    - name: varlog
      hostPath:
        path: /var/log
    - name: varlibdockercontainers
      hostPath:
        path: /var/lib/docker/containers
```

DaemonSets

Problem: Run one (and only one) pod on every node (or a subset of nodes)

Today:

- Over-replicate
- Static pods
- Grab unique resource per node
- Run outside of Kubernetes

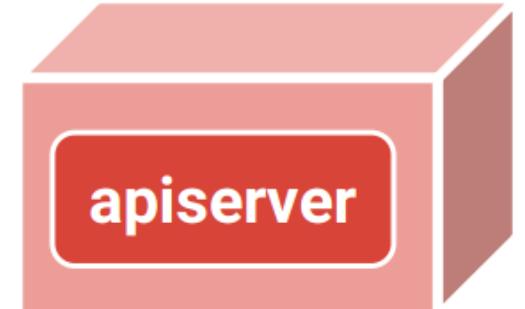


DaemonSets

Solution: Daemon Sets!

```
$ kubectl create -f daemon-sharder.yml
```

```
apiVersion: extensions/v1
kind: DaemonSet
spec:
  template:
    spec:
      nodeSelector:
        app: datastore-node
      containers:
        image: kubernetes/sharded
        ports:
          - containerPort: 9042
```

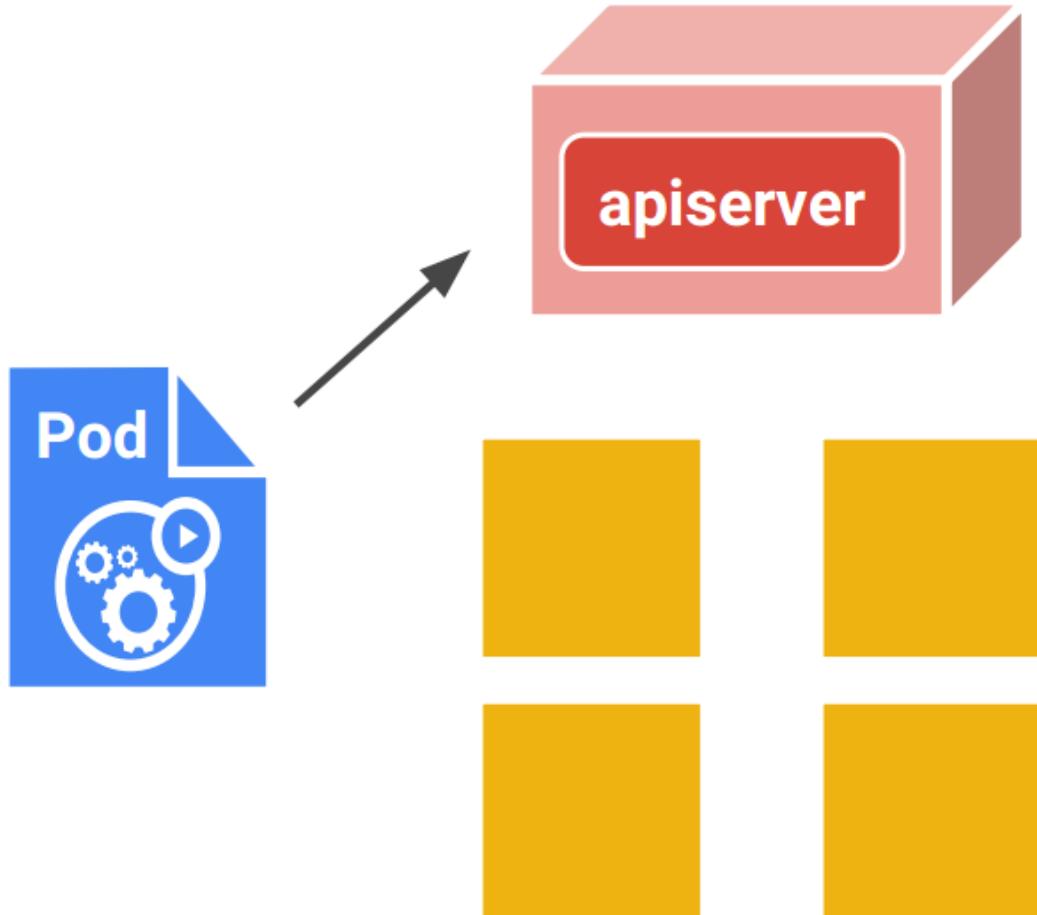


DaemonSets

Solution: Daemon Sets!

```
$ kubectl create -f daemon-sharder.yml
```

```
apiVersion: extensions/v1
kind: DaemonSet
spec:
  template:
    spec:
      nodeSelector:
        app: datastore-node
      containers:
        image: kubernetes/sharded
        ports:
          - containerPort: 9042
```

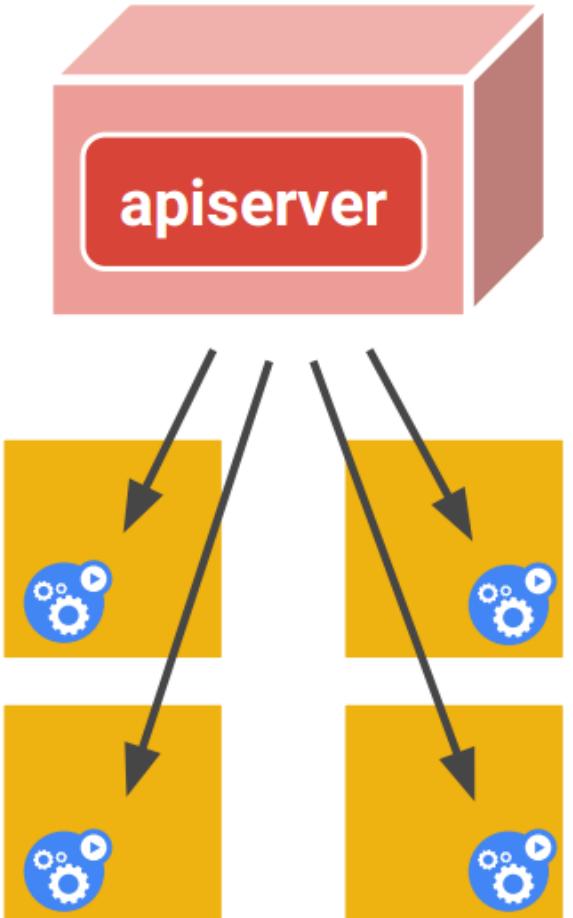


DaemonSets

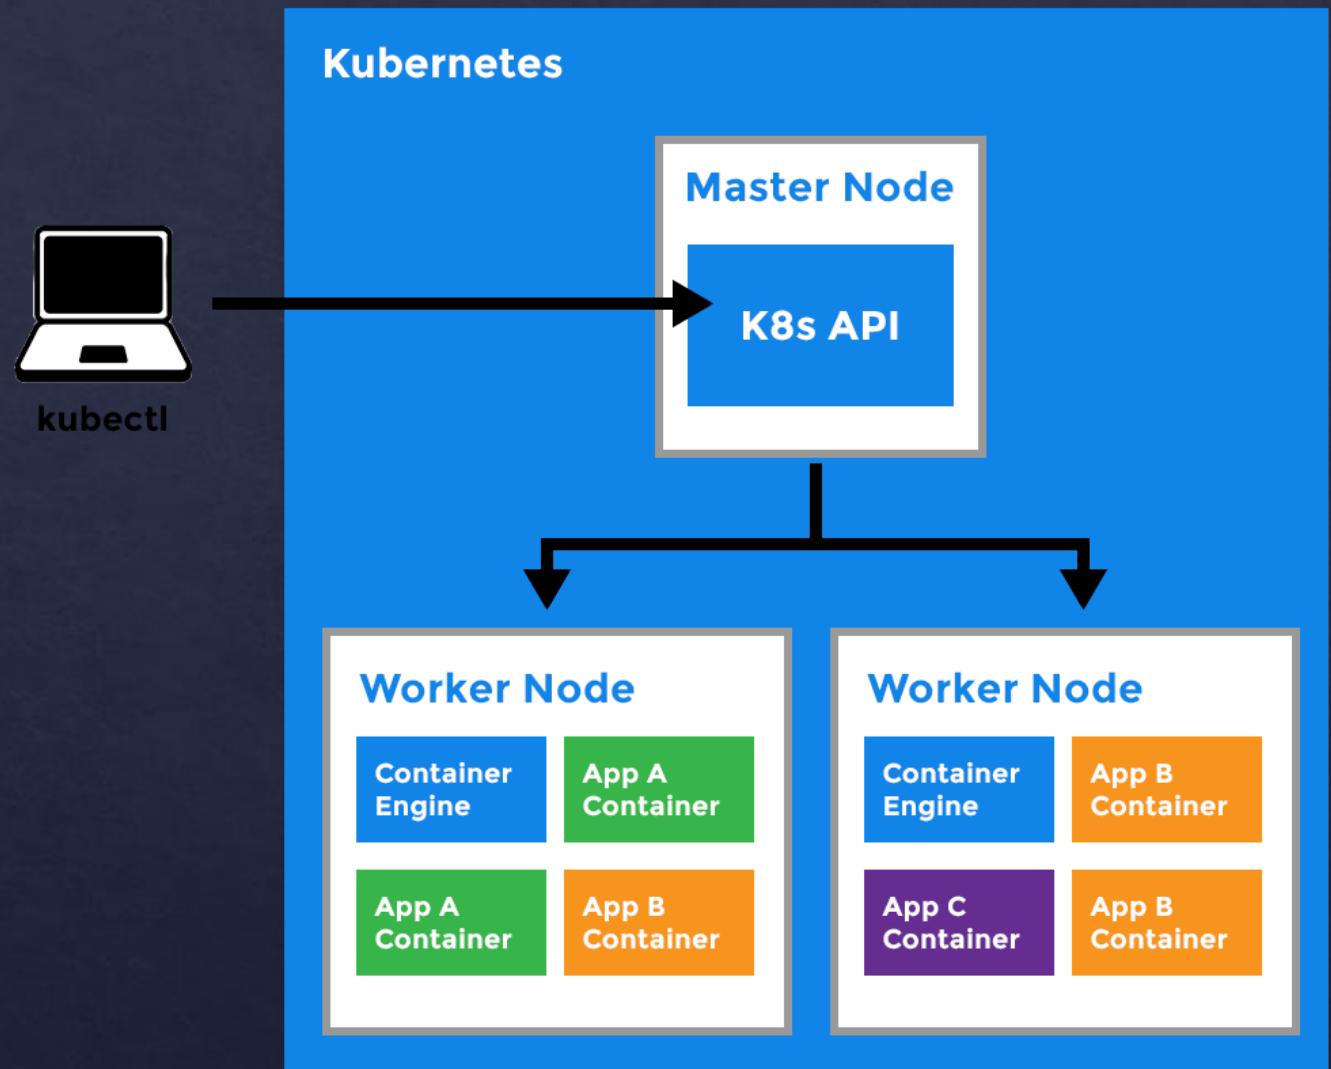
Solution: Daemon Sets!

```
$ kubectl create -f daemon-sharder.yml
```

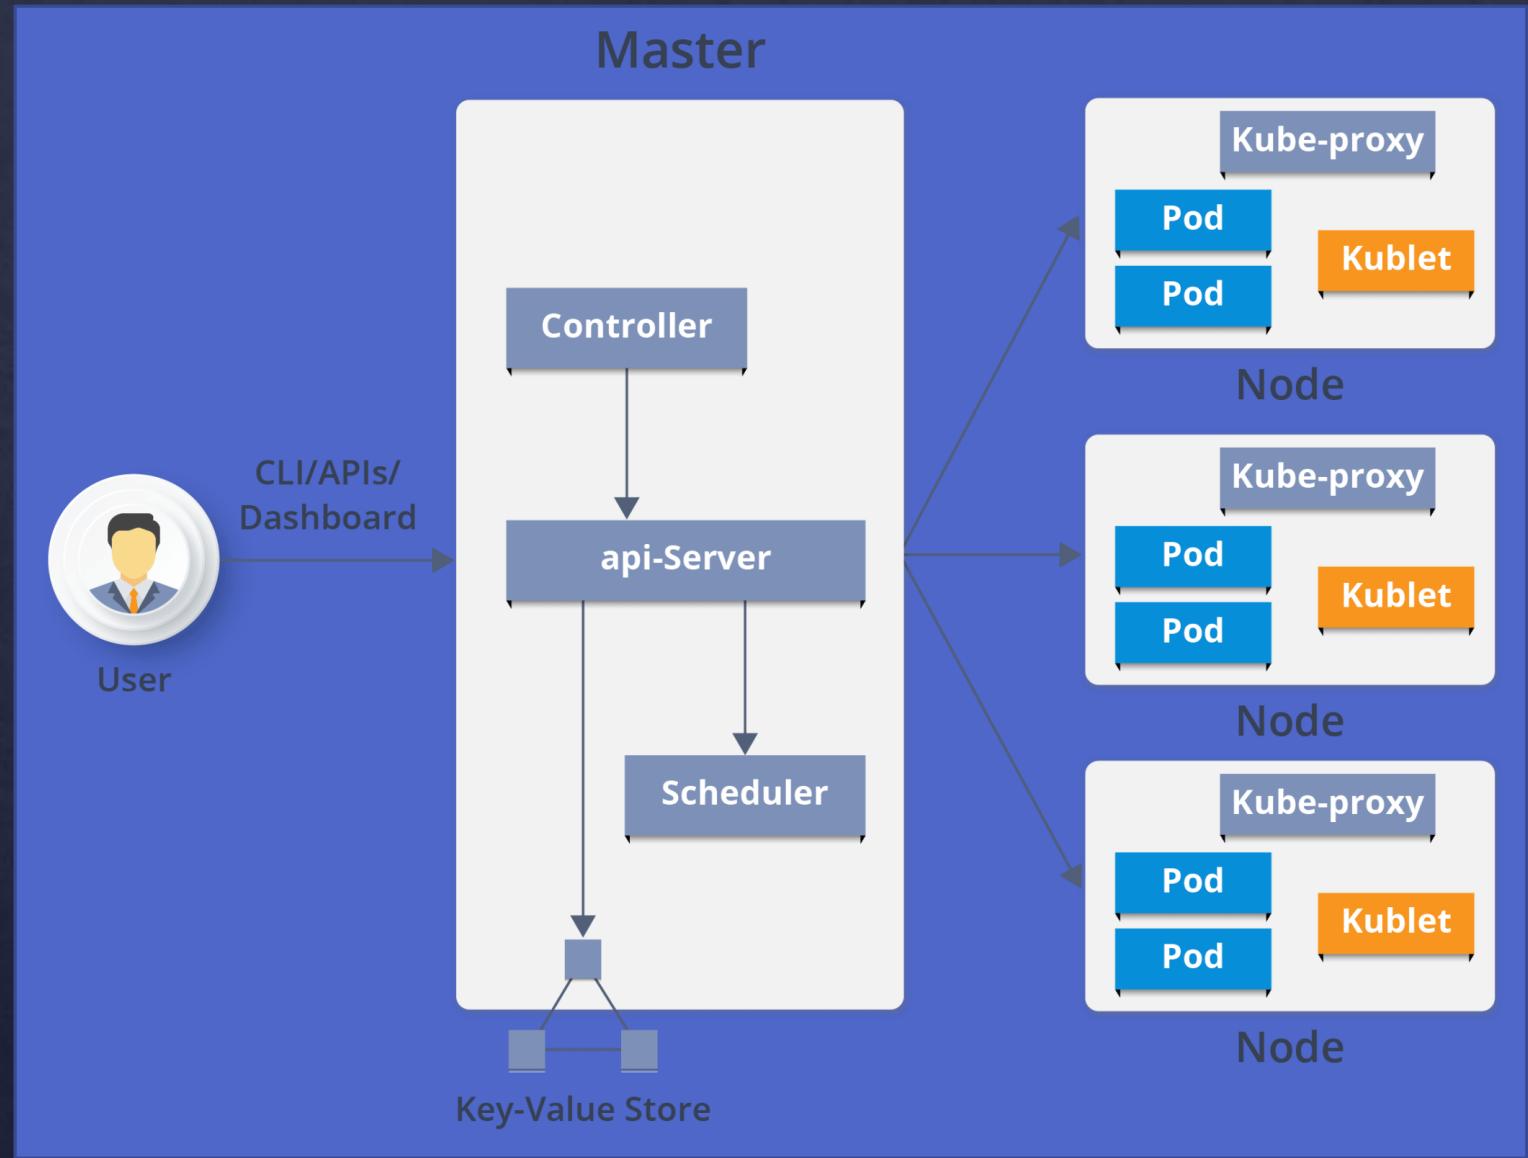
```
apiVersion: extensions/v1
kind: DaemonSet
spec:
  template:
    spec:
      nodeSelector:
        app: datastore-node
      containers:
        image: kubernetes/sharded
        ports:
          - containerPort: 9042
```



a CLI tool for Kubernetes



To be continued...



What is Kubernetes?

At its basic level: is a

...1... for ...2..., ...3... across ...4...

What is Kubernetes?

At its basic level: is a

...1... for ...2..., ...3... across ...4...

1. System

What is Kubernetes?

At its basic level: is a

System for ...**2**... , ...**3**... across ...**4**...

1. System
2. Running and Coordinating

What is Kubernetes?

At its basic level: is a

System for Running and Coordinating , ...3... across ...4...

1. System
2. Running and Coordinating
3. Containerized Applications

What is Kubernetes?

At its basic level: is a
System for Running and Coordinating , Containerized
Applications across ...**4**...

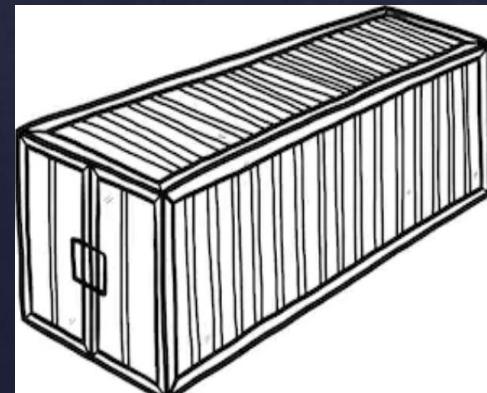
1. System
2. Running and Coordinating
3. Containerized Applications
4. Clusters of Machines

What is Kubernetes?

At its basic level: is a ...

System for Running and Coordinating , Containerized

Applications across Clusters of Machines



UI

The screenshot shows the Kubernetes Dashboard interface for deploying a containerized application. The title bar says "Kubernetes Dashboard". The main header is "kubernetes". The main content area is titled "Deploy a Containerized App".
The first section has two radio button options:

- Specify app details below
- Upload a YAML or JSON file

To learn more, [take the Dashboard Tour](#).

App name: `nginx-app` (9/24)
Container image: `nginx`
Number of pods: `12`
Port: `80` Target port: `TCP`
Protocol: `TCP`
Expose service externally:
SHOW ADVANCED OPTIONS
DEPLOY CANCEL

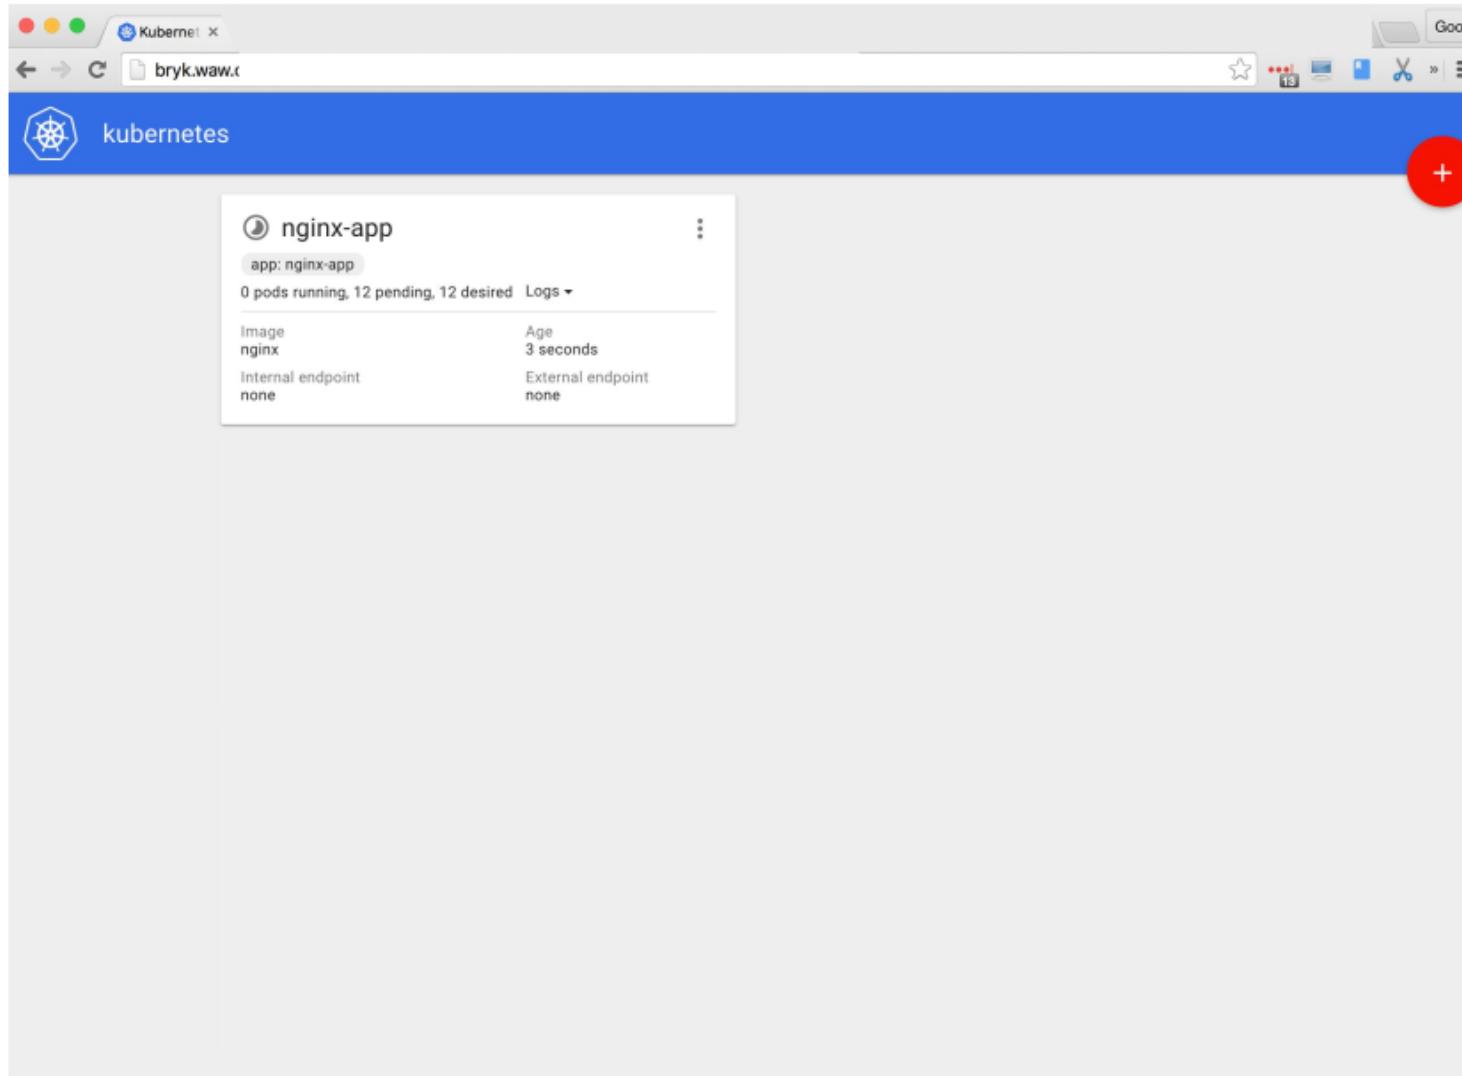
An 'app' label with this value will be added to the Replication Controller and Service that get deployed. [Learn more](#)

Enter the URL of a public image on any registry, or a private image hosted on Docker Hub or Google Container Registry. [Learn more](#)

A Replication Controller will be created to maintain the desired number of pods across your cluster. [Learn more](#)

Ports are optional. If specified, a Service will be created mapping the Port (incoming) to a target Port seen by the container. The internal DNS name for this Service will be: `nginx-app`. [Learn more](#)

UI



UI

Kubernetes Dashboard × Goog

← heapster-v11

PODS EVENTS DETAILS

↑ Pod	Status ⓘ	Restarts ⓘ	Age ⓘ	CPU (cores) ⓘ	Memory (B) ⓘ	Cluster IP ⓘ	Node ⓘ	Logs
heapster-v11-xaum1	Running	0	13 days	0.008	39.309 Mi	10.164.4.17	gke-bryk-guestbook-75b28419-node-dos0	Logs ↗

Activities Terminal



kubernetes-u18@ubuntu:~



File Edit View Search Terminal Help

kubernetes-u18@ubuntu:~\$



❖ Minikube/Dashboard

Velocity

5k Commits
in 1.2

+50% Unique
Contributors

Top 0.01% of
all Github
Projects

1200+ External
Projects Based
on K8s

Companies Contributing



Companies Using



Resources

- <https://slideplayer.com/slide/7483962/>
- <https://www.slideshare.net/chhattanshah/cluster-and-grid-computing>
- <https://martinfowler.com/articles/microservices.html>
- <http://energystorage.org/membership/members>
- <https://www.slideshare.net/Docker/introduction-to-docker-2017>
- <https://www.cloudbees.com/blog/clustering-jenkins-kubernetes-google-container-engine>

Resources

- https://en.wikipedia.org/wiki/Representational_state_transfer#History
- <http://www.studytrails.com/devops/docker-architecture-engine-containerd-runc/>
- <https://dev.to/jibinliu/how-to-persist-data-in-docker-container-2m72>
- <https://docs.docker.com/storage/volumes/>
- <https://www.slideshare.net/GiacomoVacca/docker-from-scratch>
- <https://www.ianlewis.org/en/how-kubeadm-initializes-your-kubernetes-master>
- <https://thenewstack.io/implementing-advanced-scheduling-techniques-with-kubernetes/>
- <https://medium.com/jorgeacetoz/kubernetes-node-components-service-proxy-kubelet-and-cadvisor-dcc6928ef58c>

Resources

- <https://blog.risingstack.com/what-is-kubernetes-how-to-get-started/>
- <https://linuxacademy.com/blog/containers/history-of-container-technology/>
- <https://blog.risingstack.com/what-is-kubernetes-how-to-get-started/>
- <https://coreos.com/kubernetes/docs/latest/replication-controller.html>
- <https://kubernetes.io/docs/concepts/workloads/controllers/deployment/>
- <https://www.slideshare.net/kubecon/kubecon-eu-2016-keynote-kubernetes-state-of-the-union>
- [DevOps With Kubernetes Book](#)